

TILL LOHFINK & STEFAN ZAPF

CORE ML BOOTCAMP: BUILD IOS & MAC APPS WITH DEEP LEARNING

Use Core ML to add intelligence to your application

PART 1

Our Bootcamp

What can AI do for you?

The Create ML Tool

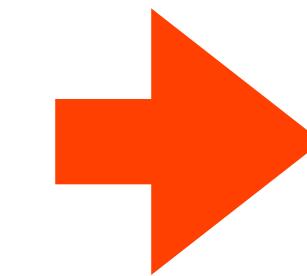
CoreML Conversion

CORE ML BOOTCAMP

**DEEP LEARNING IS MACHINE LEARNING WITH
AUTOMATED FEATURE ENGINEERING**

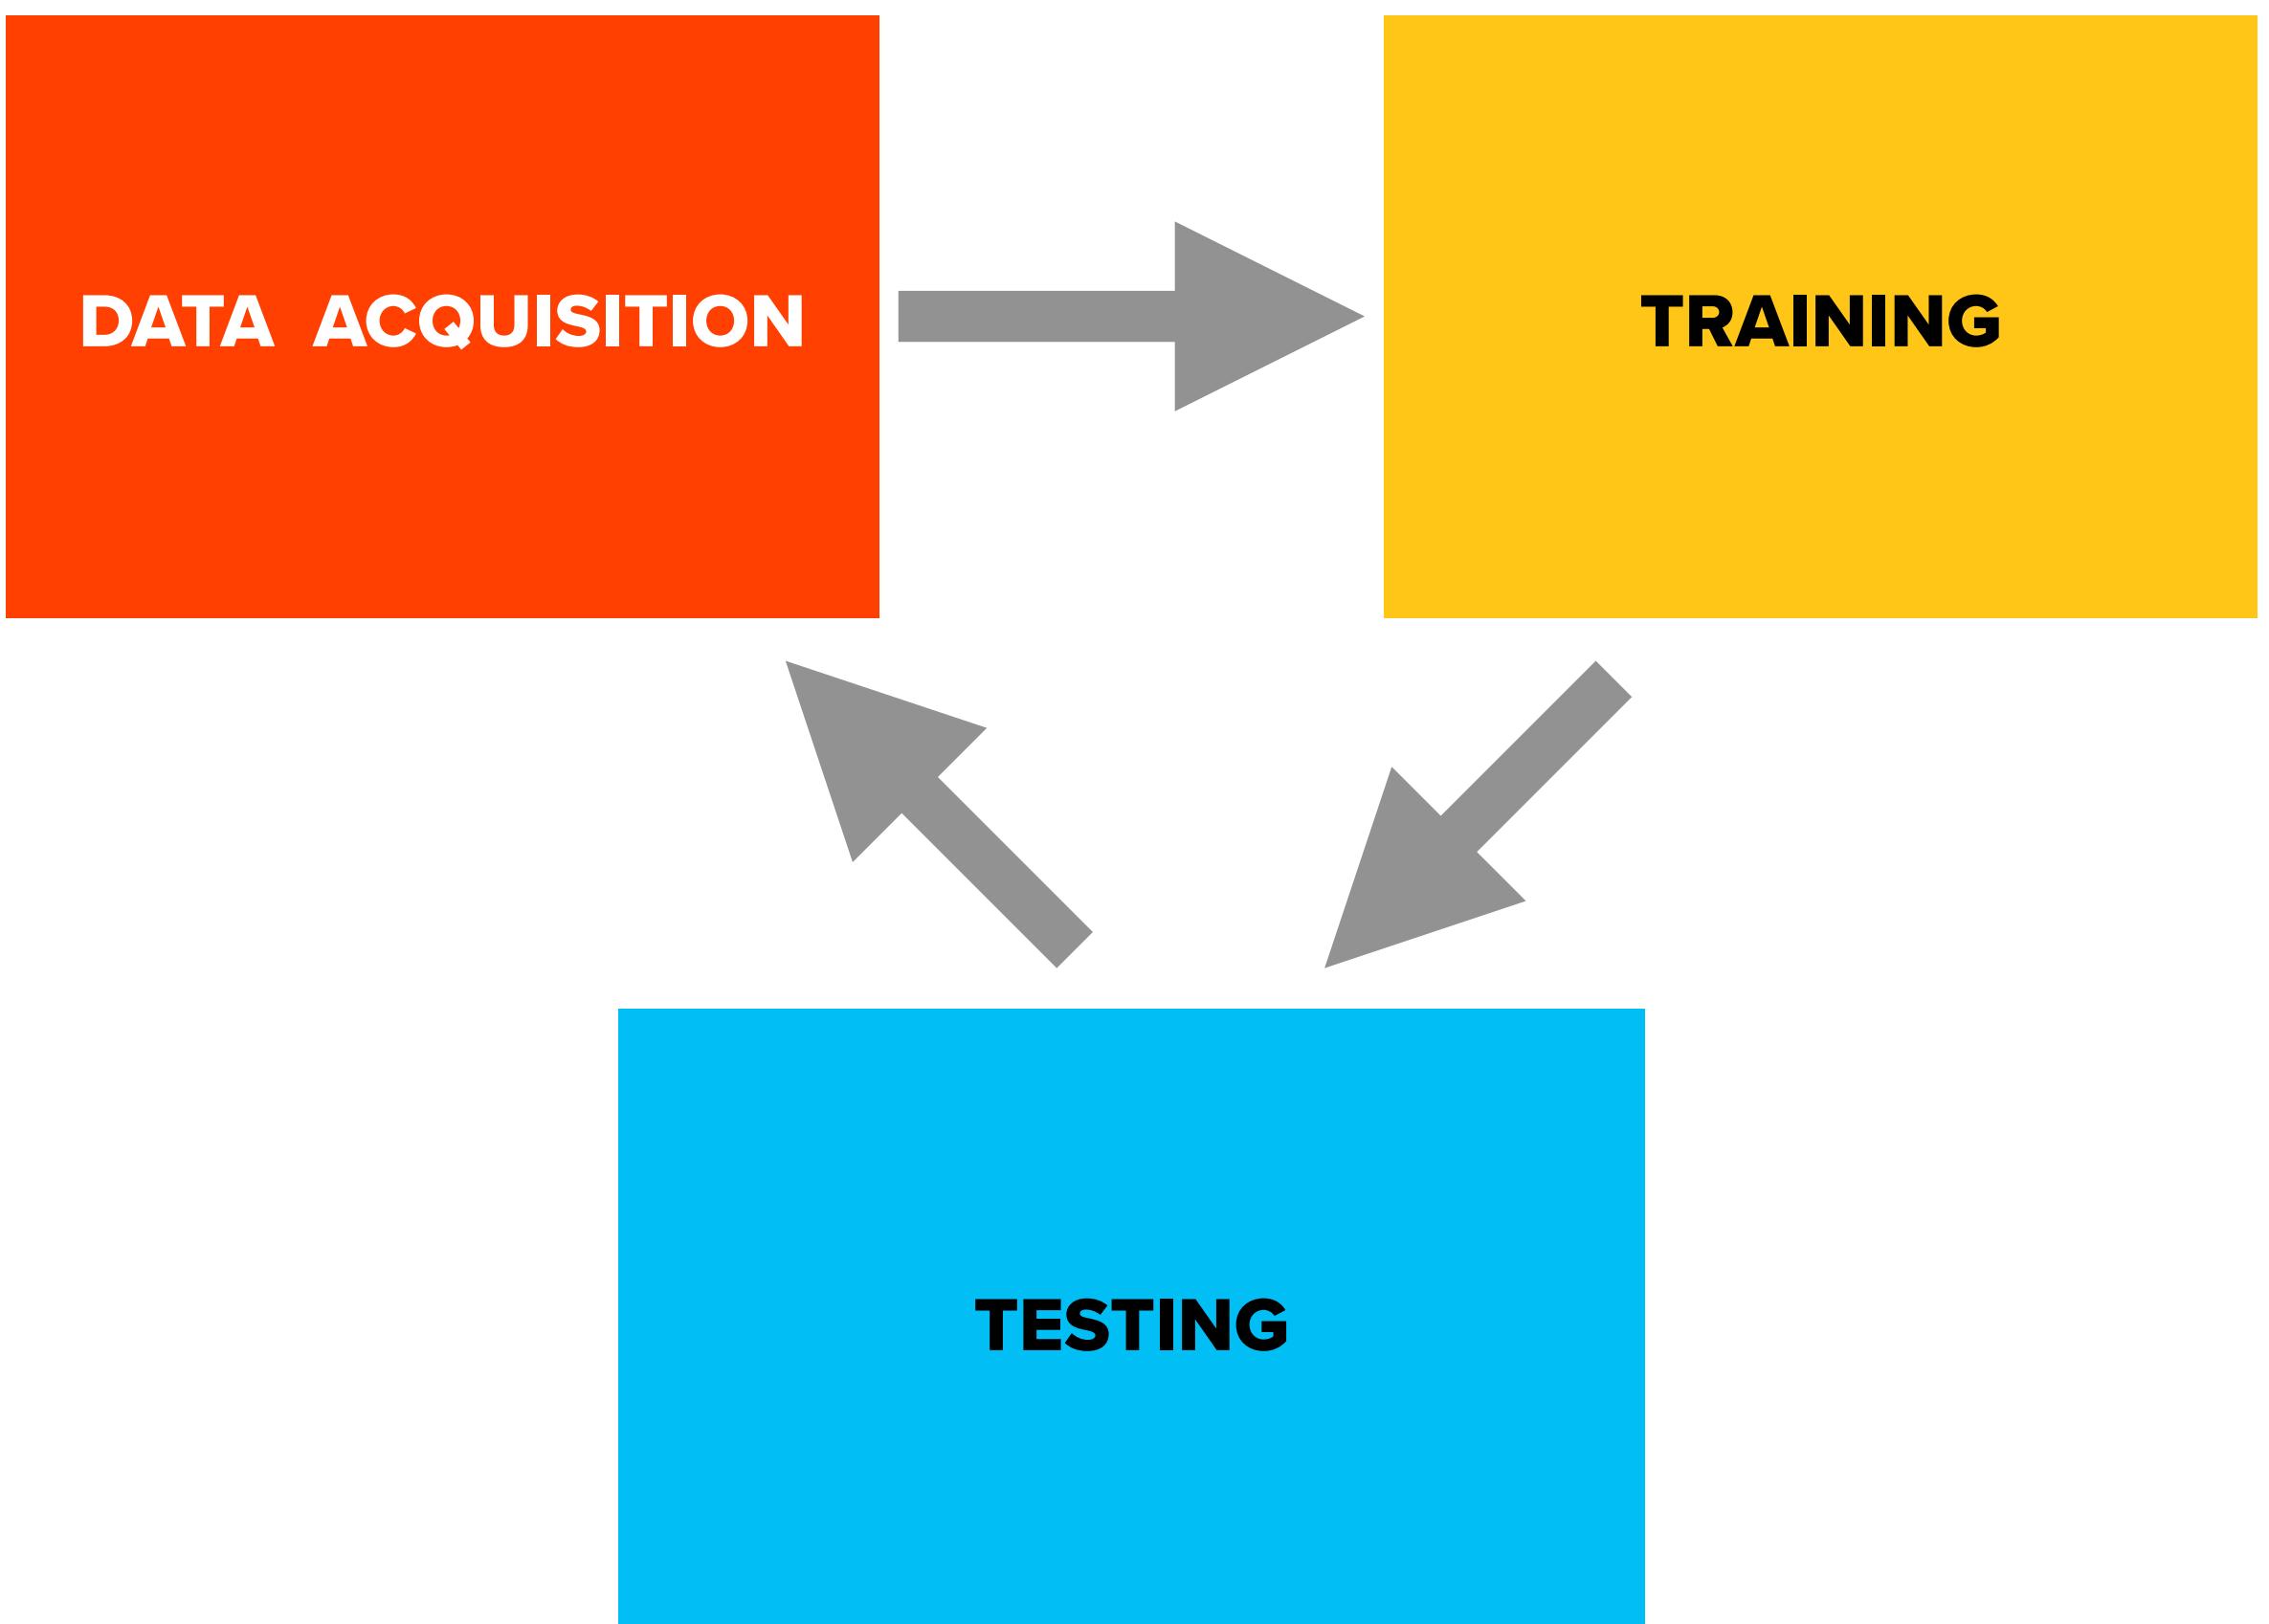
BEFORE & AFTER DL1

- Before Picture
- Object Detection:
 - Edge Detection
 - Template Matching
- Each method had problems:
 - Lighting Conditions
 - Different Sizes / Angles
 - Weather Conditions such as snow / rain ...
- A lot of research on all the sub problems
- Careers were mad in researching these sub problems



BEFORE & AFTER DL1

- Today with Deep Learning
- Data, Training, Testing - Triad
- If lightning is a problem
 - Find new data with different lightning
 - Train again
 - Test again
 - Repeat



FOUR LEVELS OF AI ENGINEERING

FOUR LEVELS OF AI ENGINEERING

- **1 Using public APIs**
 - e. g. Azure Cognitive Services, OpenAI, etc.
- **2 Utilize Transfer Learning with existing state-of-the-art solutions**
 - Like in this class
 - Build on top of leading architectures like BERT or ResNet
- **3 Coming up with your own models**
 - Work directly in PyTorch or Tensorflow
 - Integrate State-Of-The-Art research
- **4 Auto ML**
 - Create AIs that can create AIs
 - Deep Reinforcement Learning

AI TASKS: WHAT CAN AI DO FOR YOU?

WHAT CAN AI DO FOR YOU?

- Computer Vision
 - Semantic Segmentation
 - Image classification
 - Object Detection
 - Image Generation
 - Pose Estimation
 - Super-Resolution
 - Object Tracking
 - Activity Recognition
- Natural Language Processing
 - Machine Translation
 - Language Modeling
 - Question Answering
 - Sentiment Analysis
 - Text Generation
 - Named Entity Recognition
 - Text Summarization
 - Semantic Textual Similarity
 - Part-Of-Speech Tagging

SOLVE THE MOST DIFFICULT PROBLEMS TO DATE WITH EASE

The screenshot shows the homepage of paperswithcode.com. At the top, there's a search bar with the placeholder "Search for paper" and a navigation bar with links for "Browse State-of-the-Art", "Methods", "Reproducibility", "Portals", and "About". There are also social media icons for Twitter and GitHub, and a "Log In/Register" button.

The main content area is titled "Trending Research" and features three cards for trending papers:

- Training data-efficient image transformers & distillation through attention**
23 Dec 2020 • facebookresearch/deit • PyTorch
We show the interest of this token-based distillation, especially when using a convnet as a teacher.
Ranked #2 on [Image Classification on iNaturalist 2018](#)
FINE-GRAINED IMAGE CLASSIFICATION
★ 807 stars / hour
- YolactEdge: Real-time Instance Segmentation on the Edge**
(Jetson AGX Xavier: 30 FPS, RTX 2080 Ti: 170 FPS)
22 Dec 2020 • haotian-liu/yolact_edge • PyTorch
We propose YolactEdge, the first competitive instance segmentation approach that runs on small edge devices at real-time speeds.
REAL-TIME INSTANCE SEGMENTATION SEMANTIC SEGMENTATION
★ 465 stars / hour
- Mastering Atari, Go, Chess and Shogi by Planning with a Learned Model**
19 Nov 2019 • werner-duvaud/muzero-general • PyTorch
1.27 stars / hour

SOLVE THE MOST DIFFICULT PROBLEMS TO DATE WITH EASE

The screenshot shows the homepage of paperswithcode.com. At the top, there's a navigation bar with links for 'Browse State-of-the-Art', 'Methods', 'Reproducibility', 'Portals', 'About', and 'Log In/Register'. Below the navigation is a large header section with a purple-to-blue gradient background. The text 'Browse State-of-the-Art' is centered, followed by a subtext '3,661 benchmarks • 1,927 tasks • 3,183 datasets • 39,045 papers with code'. There are social media links for Twitter and GitHub, and a 'Follow on Twitter for updates' button.

Computer Vision

- Semantic Segmentation**: 79 benchmarks, 1431 papers with code
- Image Classification**: 175 benchmarks, 1250 papers with code
- Object Detection**: 147 benchmarks, 1051 papers with code
- Image Generation**: 131 benchmarks, 501 papers with code
- Pose Estimation**: 229 benchmarks, 461 papers with code

[See all 958 tasks](#)

Natural Language Processing

- Machine Translation**: 56 benchmarks, 971 papers with code
- Language Modelling**: 16 benchmarks, 946 papers with code
- Question Answering**: 65 benchmarks, 847 papers with code
- Sentiment Analysis**: 50 benchmarks, 579 papers with code
- Text Generation**: 47 benchmarks, 407 papers with code

[See all 350 tasks](#)

**HOW CAN YOU USE THESE AI CAPABILITIES IN
YOUR APP?**

DEMO TIME: CREATE ML TOOL

VISION FRAMEWORK

- Out-Of-The-Box components to build image and video processing applications
- Quite extensive; offers solutions for image classification, similarity analysis and object detection with little code
- Great for moving objects featuring object tracking and even trajectory detection
- Neat stuff like text detection, lots of health related algorithms like pose detection etc.

- In general works like this: you bring an image and pass it to an instance of an `VNImageRequestHandler` which can call an instance of the desired `VNSomethingSomethingRequest`
- We will use this framework later!

NATURAL LANGUAGE FRAMEWORK

- Out-Of-The-Box components to build text-based applications - to be more specific
- Not as extensive as Apple's Vision Framework
- Brings some great tools for tokenization and lemmatization (finding a word's stem; much more interesting in romance languages)
- Can help you with a lot of tagging feature like named entity recognition or part-of-speech tagging
- Very easy to implement in your application: Instantiate the desired class and call a method passing your text
- If you want to do more sophisticated things like sentiment analysis, Q & As etc. you have to bring your own models / implementations (i. e. Papers With Code)

WHAT CAN APPLE DO FOR YOU? (OUT OF THE BOX)

- Computer Vision
 - Semantic Segmentation
 - Image classification ✓
 - Object Detection ✓
 - Image Generation
 - Image Similarity Analysis ✓
 - Pose Estimation ✓
 - Super-Resolution
 - Object Tracking ✓
 - Activity Recognition ✓
- Natural Language Processing
 - Machine Translation
 - Language Modeling
 - Question Answering
 - Sentiment Analysis
 - Text Generation
 - Named Entity Recognition
 - Text Summarization
 - Semantic Textual Similarity
 - Part-Of-Speech Tagging

WHAT CAN APPLE DO FOR YOU? (OUT OF THE BOX)

- Computer Vision
 - Semantic Segmentation
 - Image classification ✓
 - Object Detection ✓
 - Image Generation
 - Image Similarity Analysis ✓
 - Pose Estimation ✓
 - Super-Resolution
 - Object Tracking ✓
 - Activity Recognition ✓
- Natural Language Processing
 - Machine Translation
 - Language Modeling
 - Question Answering
 - Sentiment Analysis ✓
 - Text Generation
 - Named Entity Recognition ✓
 - Text Summarization
 - Semantic Textual Similarity
 - Part-Of-Speech Tagging ✓

IMPORTING EXISTING MODELS OUTSIDE OF APPLE'S ECOSYSTEM

HUGGINGFACE

DEMO TIME: GOOGLE COLAB

INTERMISSION

PART 2: THE LEGAL APP

App Specifications

The Transformer & Bert

Coding 1: The UI

Coding 2: The Pdf Component

Coding 3: Using BERT with GCD

THE LEGAL APP

LEGAL APP PURPOSE

- A legal argument is based on:
 - facts, and
 - legal argument, mostly applying or distinguishing precedent
- When presenting a case, the lawyer will need to find the facts in a large corpus of documents that support their legal argument and build defenses against facts that do not confirm the legal narrative
- During discovery, the pertinent fact can be buried in a large corpus of documents
- Large Law firms have a staff of paralegals and other support staff to find the needle in a haystack
- Smaller law firms are more affordable, but often don't have the staff to tackle a large corpus
- Lawyers can ask questions about facts to the legal app to find the right facts and thereby leveling the playing field

YOUR AI HELPS DEMOCRATIZE LEGAL REPRESENTATION

THE TRANSFORMER

Attention Is All You Need

Ashish Vaswani*
Google Brain
vaswani@google.com

Noam Shazeer
Google Brain
noam@google.com

Niki Parmar*
Google Research
nikip@google.com

Llion Jones*
Google Research
llion@google.com

Aidan N. Gomez* †
University of Toronto
`aidan@cs.toronto.edu`

Łukasz Kaiser*
Google Brain
lukasz.kaiser@google.com

Illia Polosukhin* ‡
illia.polosukhin@gmail.com

ORIGIN PAPER FOR TRANSFORMERS

BERT IS A SPECIFIC TRANSFORMER

Attention is all you need

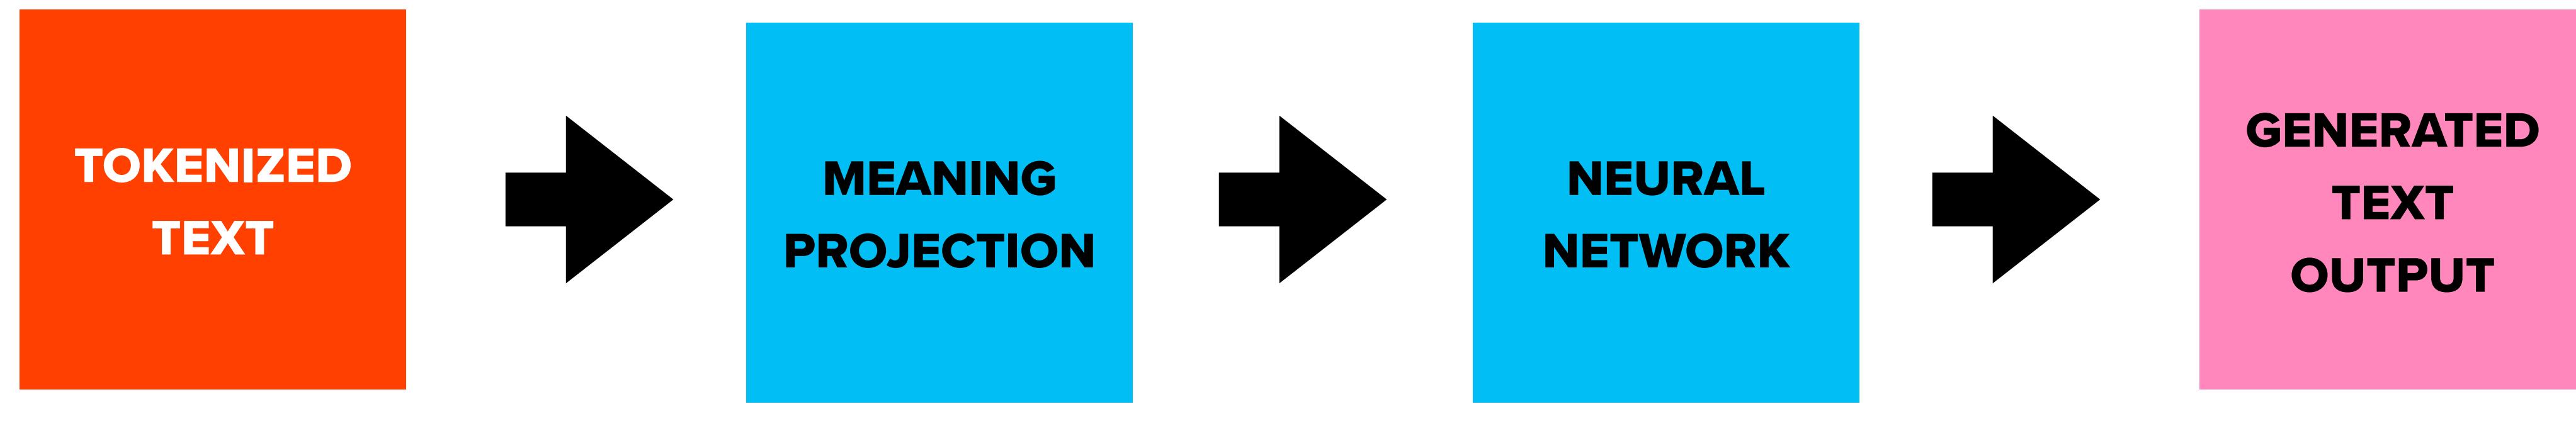
A Vaswani, N Shazeer, N Parmar... - Advances in neural ... , 2017 - papers.nips.cc

The dominant sequence transduction models are based on complex recurrent or convolutional neural networks in an encoder and decoder configuration. The best performing such models also connect the encoder and decoder through an attentionm ...

  Cited by 16137 Related articles All 20 versions 



BIRD'S EYE'S VIEW OF NATURAL LANGUAGE UNDERSTANDING



Unsupervised Pre-Trained Task

- Masked Language Modelling
- Next Sentence Prediction
- Causal or Bi-Directional

PRE-TRAINED
TO
UNDERSTAND
A SET OF
LANGUAGE(S)

Supervised Downstream Task

- Q&A
- Neural Machine Translation
- Chat Bot
- ... Anything you train it to do

THE EVOLUTION OF MEANING PROJECTION

- Word Embeddings
- Transformers

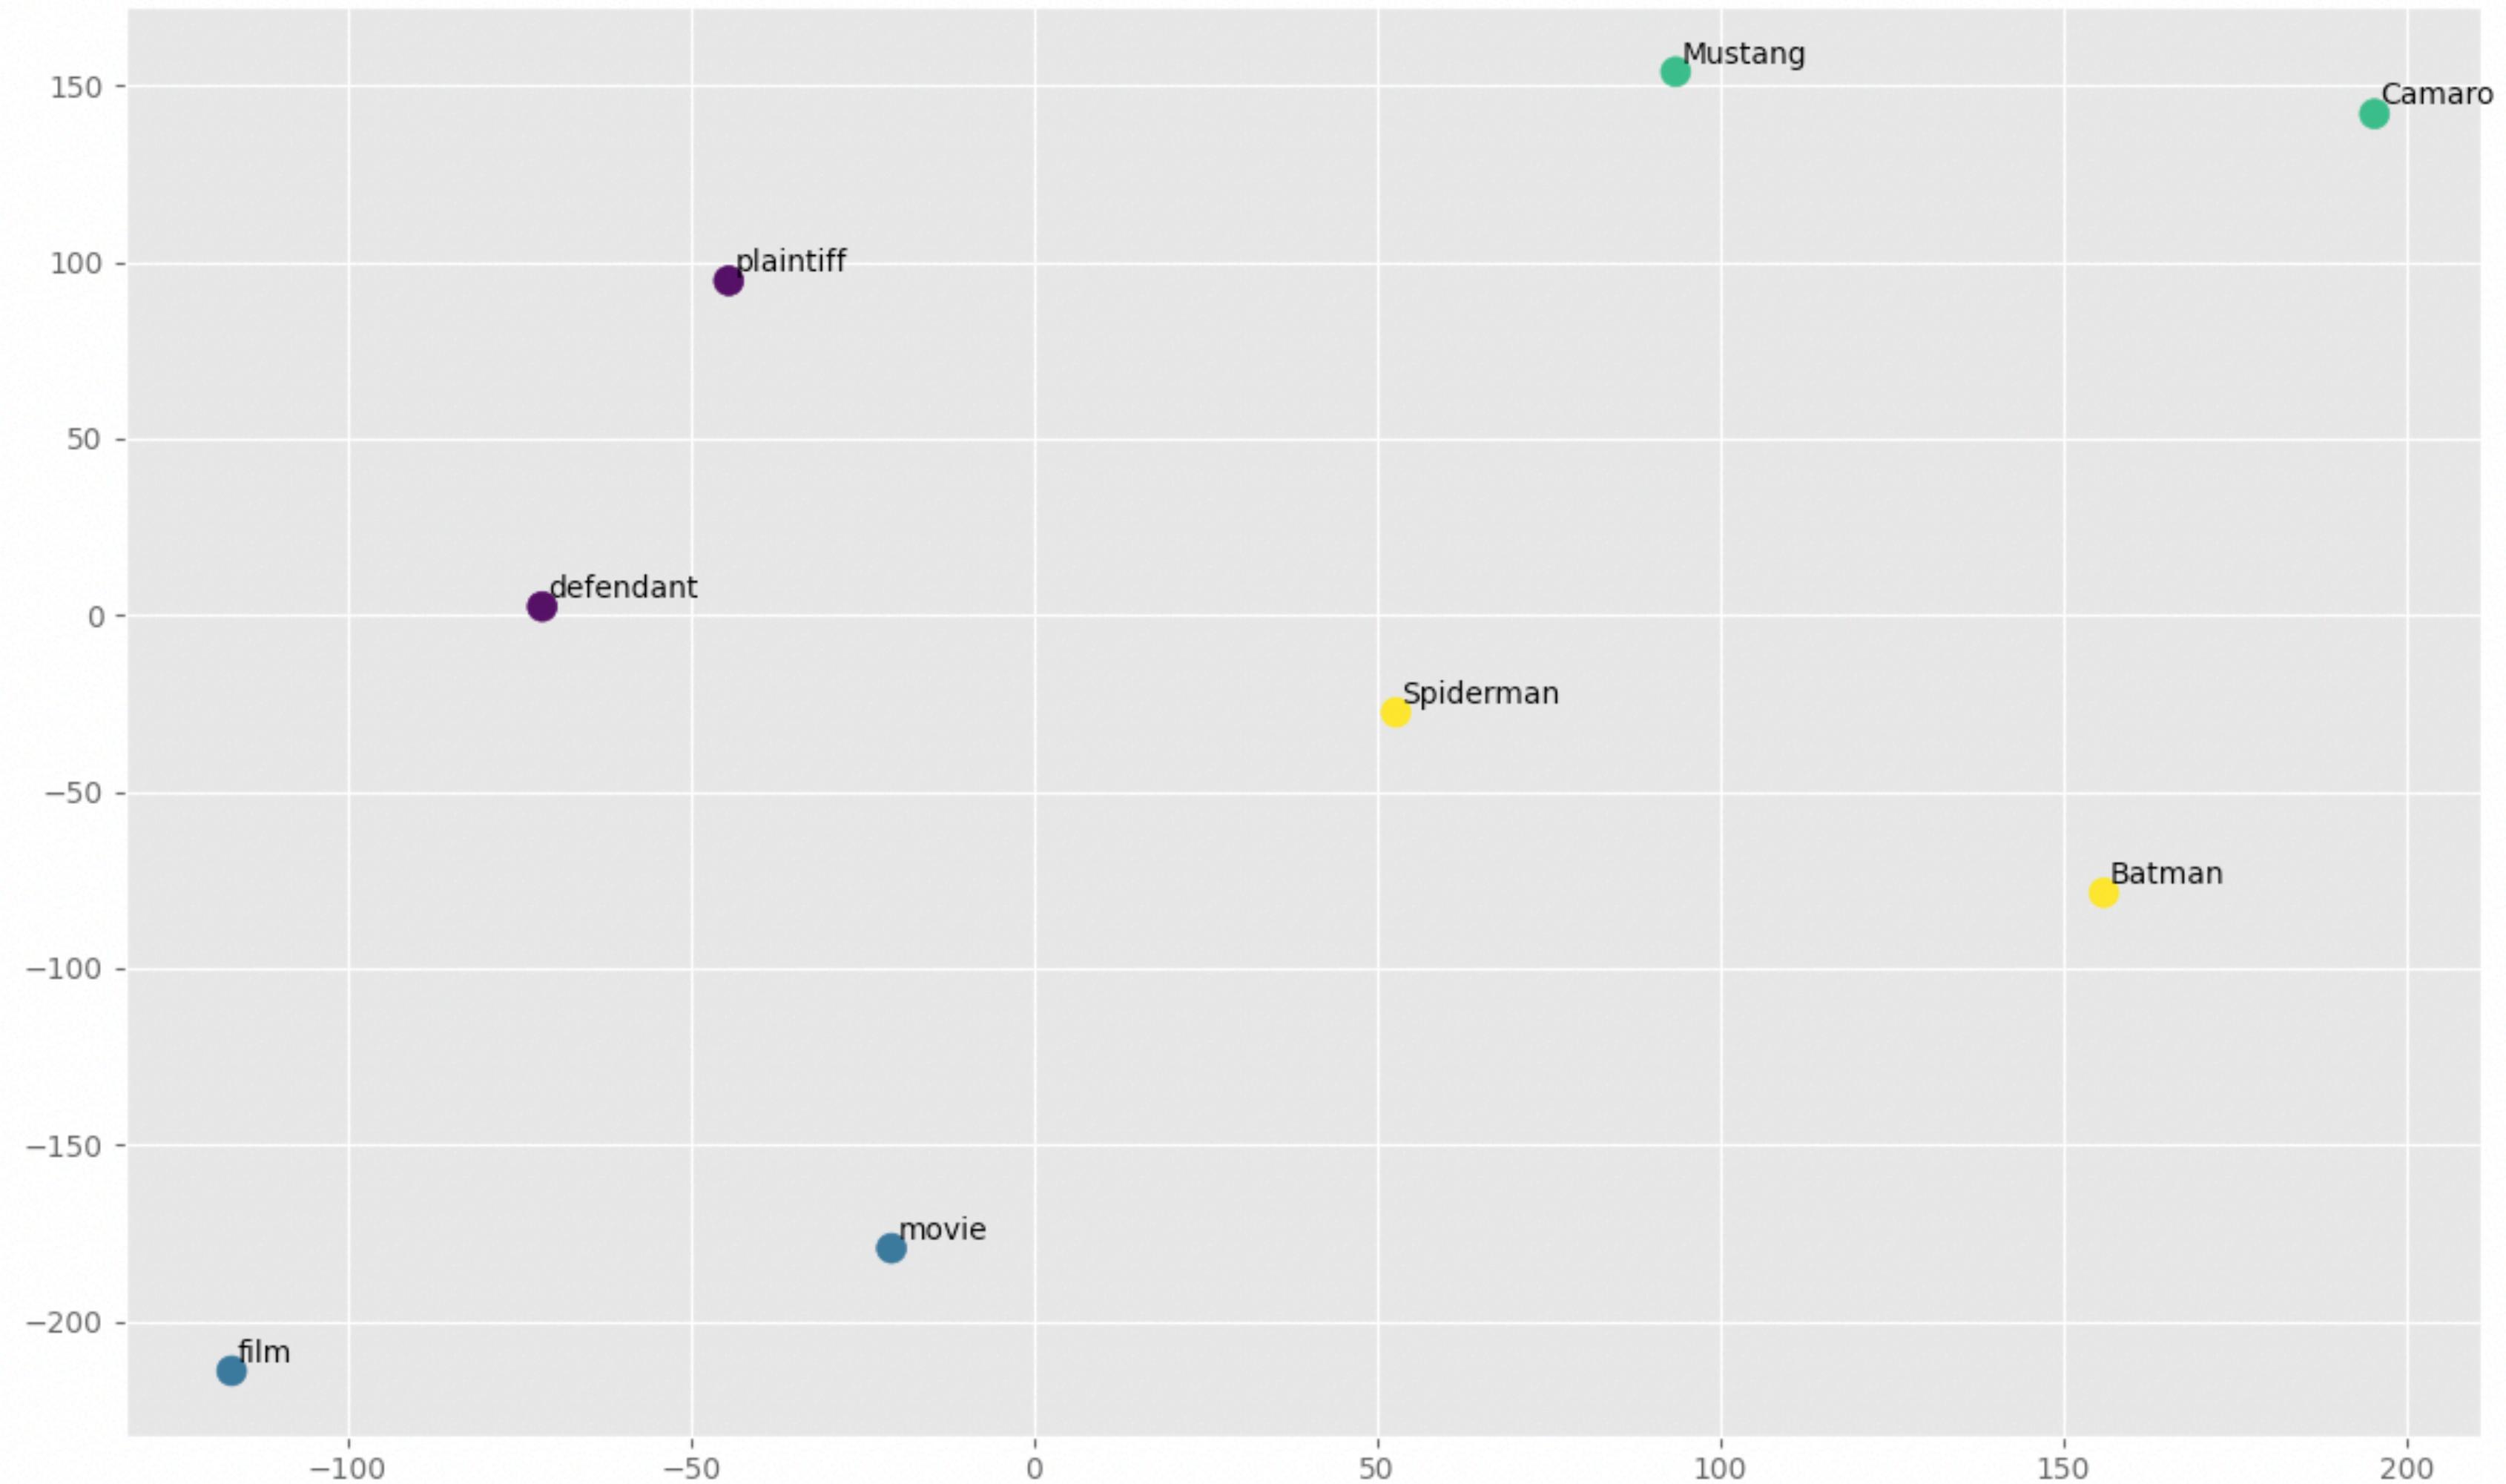
WORD EMBEDDINGS, MEANING AND LATENT SPACE

LATENT SPACE OF MEANING

EVERY WORD HAS A PLACE
IN THE LATENT SPACE

WORDS THAT HAVE A
HIGH SEMANTIC
SIMILARITY
ARE CLOSER IN THE
LATENT SPACE

METHOD:
WORD EMBEDDING
WORD -> VECTOR



**CHAT STORMING: WHAT COULD BE A PROBLEM WITH WORD
EMBEDDINGS?**

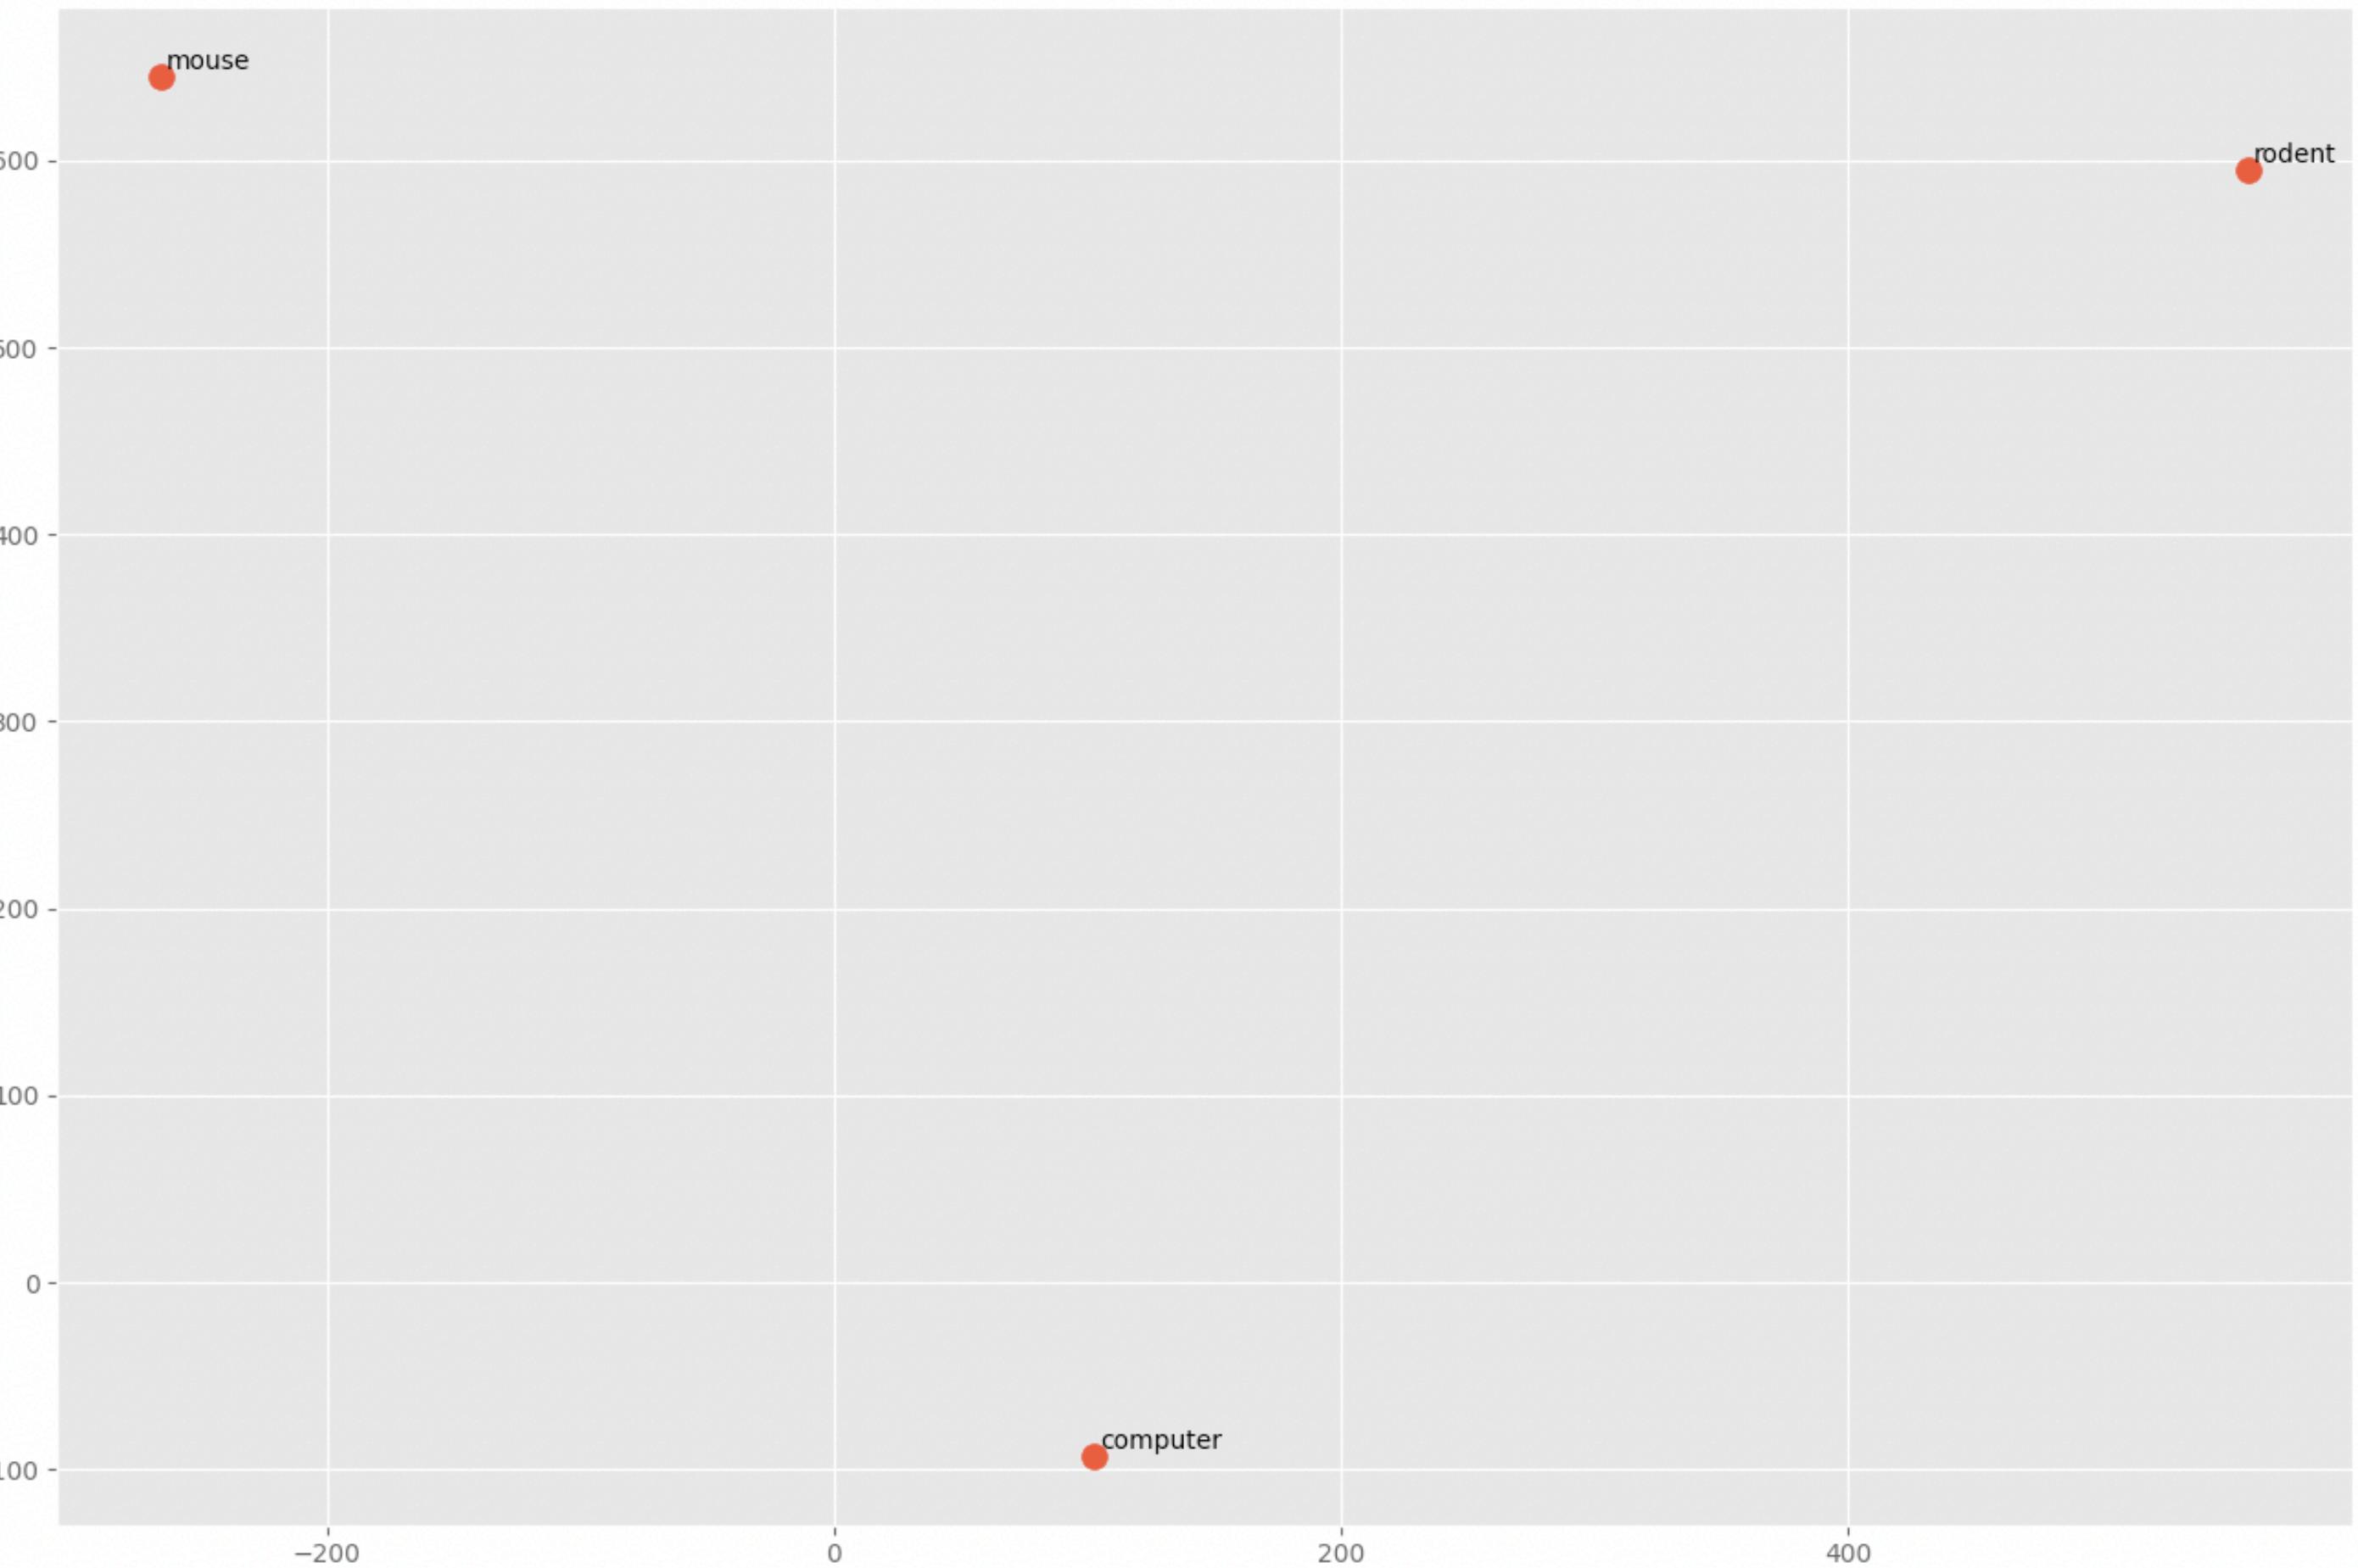
PLEASE WRITE YOUR IDEAS IN THE CHAT

AMBIGUITIES

SYMBOL “MOUSE”

POINTS TO TWO MEANINGS

**AMBIGUOUS CLOSENESS IN
LATENT SPACE OF MEANING**



**POLL: HOW CAN YOU
DISAMBIGUATE THE DIFFERENT MEANINGS
‘THE ARTISTS DRAWS A GUN’?**

AMBIGUITY OF WORDS

- The artist draws a gun.
- Two meanings
 - Artistic rendering
 - A dangerous situation

DEMO TIME: BERT DISAMBIGUATES



HUGGING FACE

[Back to all models](#)

Model: google/bert2bert_L-24_wmt_en_de

pytorch

encoder-decoder

seq2seq

en

de

dataset:wmt14

arxiv:1907.12461

license:apache-2.0

translation

[Model card](#)

[Files and versions](#)

[Use in transformers](#)

Hosted inference API ⓘ

translation

The artist draws a gun.

Compute

<s> der Künstler zieht eine Pistole." </s>

▶ JSON Output

↔ API endpoint ⓘ

Share []

DRAWS -> ZIEHT

DANGEROUS SITUATION

Contributed by



Google AI

company

6 team members · 116 models

[Model card](#)

Hosted on [huggingface.co](#)



HUGGING FACE

[Back to all models](#)

Model: `google/bert2bert_L-24_wmt_en_de`

pytorch

encoder-decoder

seq2seq

en

de

dataset:wmt14 📁

arxiv:1907.12461

license:apache-2.0

translation

Model card

Files and versions

⟨/⟩ Use in transformers

ADD: WITH A PENCIL

DRAWS -> ZEICHNET

ARTISTIC RENDERING

Hosted inference API ⓘ

translation

The artist draws a gun with a pencil.

Compute

<s> der Künstler zeichnet eine Pistole mit einem Bleistift.
</s>

▶ JSON Output

↔ API endpoint ⓘ

Share []

Contributed by



Google AI

company

6 team members · 116 models

Model card

Hosted on [huggingface.co](#)



HUGGING FACE

[⬆ Back to all models](#)

Model: `google/bert2bert_L-24_wmt_en_de`

pytorch

encoder-decoder

seq2seq

en

de

dataset:wmt14

arxiv:1907.12461

license:apache-2.0

translation

[Model card](#) [Files and versions](#)

[⟨/⟩ Use in transformers](#)

Hosted inference API ⓘ

translation

The soldier draws a gun with a pencil.

Compute

<s> der Soldat zieht ein Gewehr mit einem Bleistift. </s>

[▶ JSON Output](#) [↔ API endpoint ⓘ](#)

Share

**REPLACE: ARTIST
BY: SOLDIER**

DRAWS -> ZIEHT

DANGEROUS SITUATION

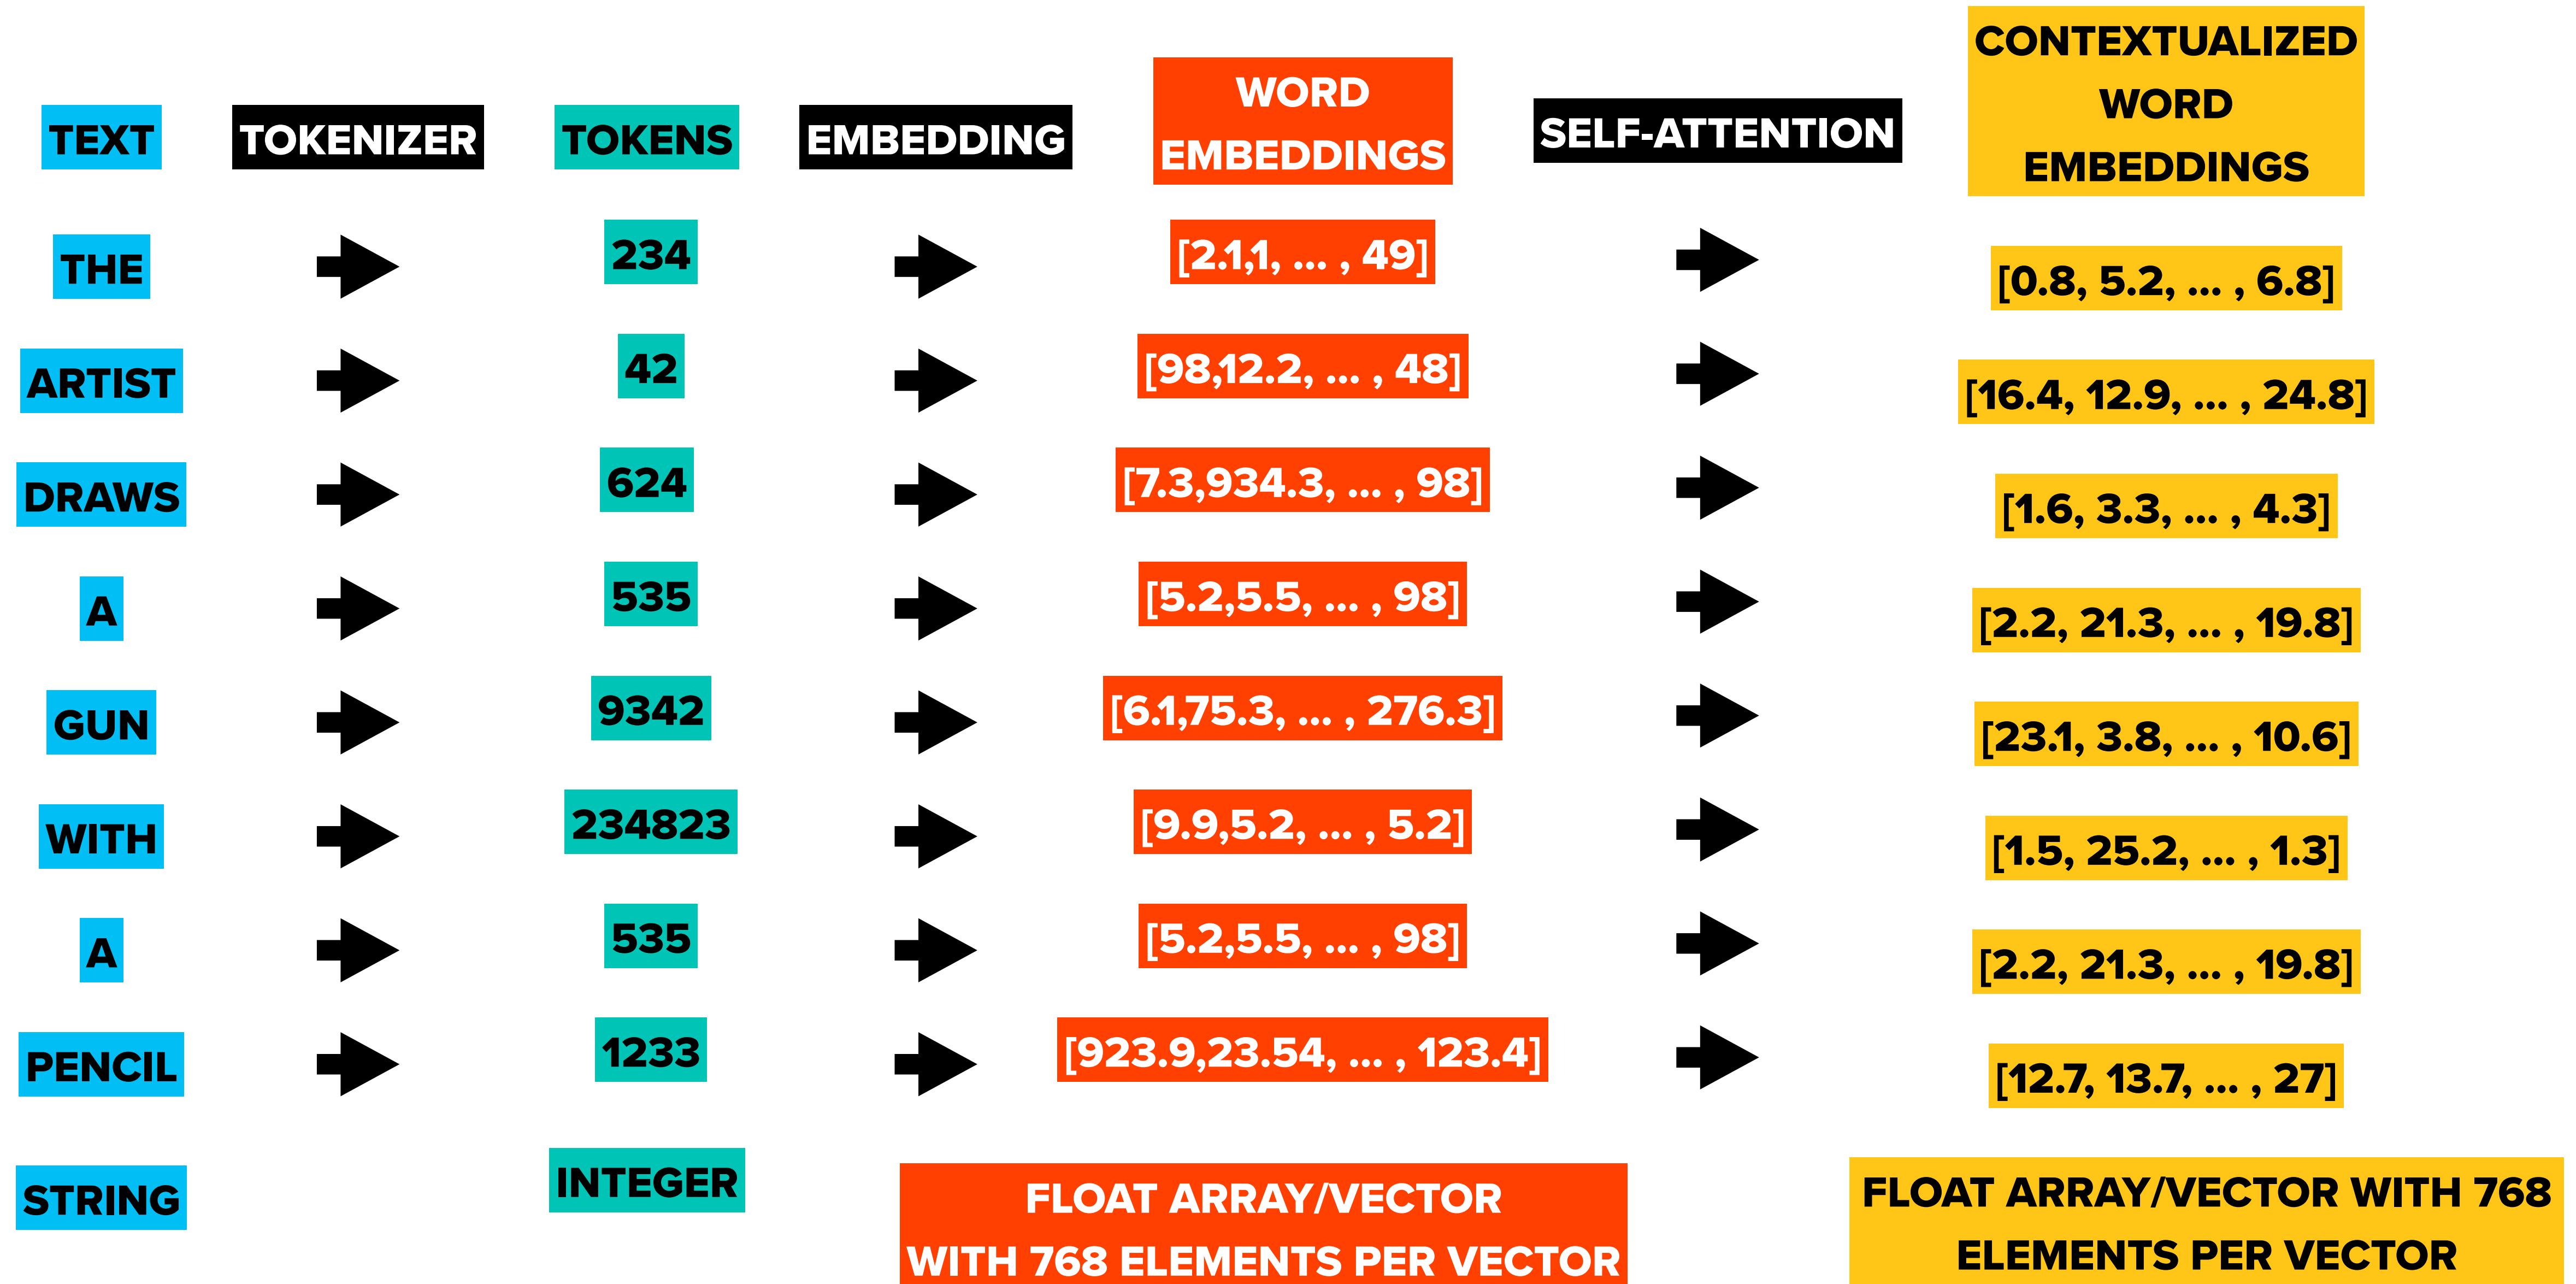
Contributed by



Google AI company

6 team members · 116 models

SELF-ATTENTION I



Attention Is All You Need

Ashish Vaswani*
Google Brain
avaswani@google.com

Noam Shazeer*
Google Brain
noam@google.com

Niki Parmar*
Google Research
nikip@google.com

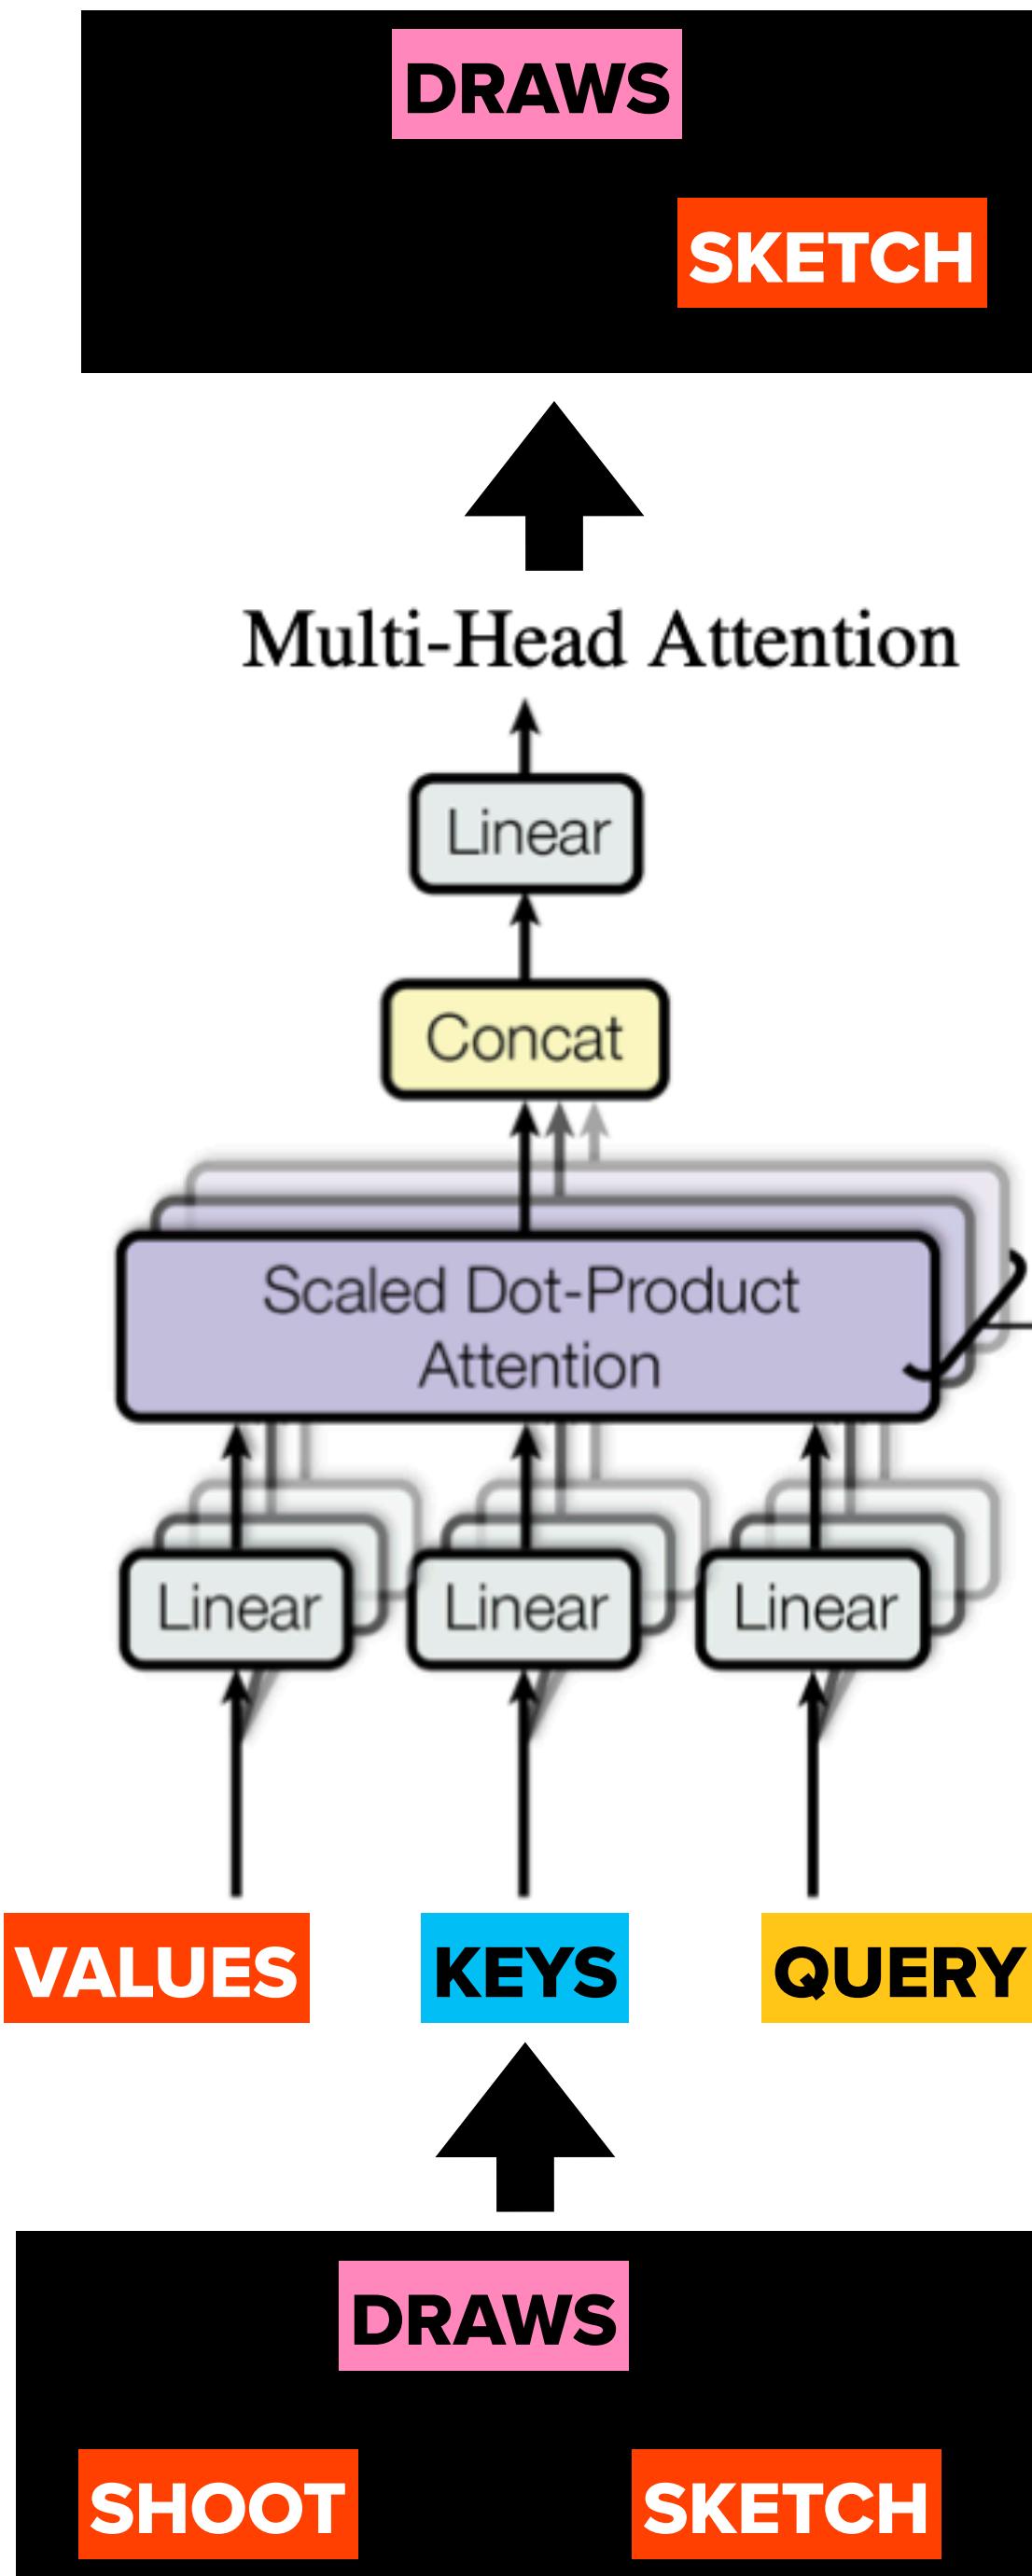
Jakob Uszkoreit*
Google Research
usz@google.com

Llion Jones*
Google Research
llion@google.com

Aidan N. Gomez* †
University of Toronto
aidan@cs.toronto.edu

Lukasz Kaiser*
Google Brain
lukasz.kaiser@google.com

Illia Polosukhin* ‡
illia.polosukhin@gmail.com



DRAWS

SKETCH

Multi-Head Attention

Linear

Concat

Scaled Dot-Product
Attention

Linear

Linear

Linear

VALUES

KEYS

QUERY

DRAWS

SHOOT

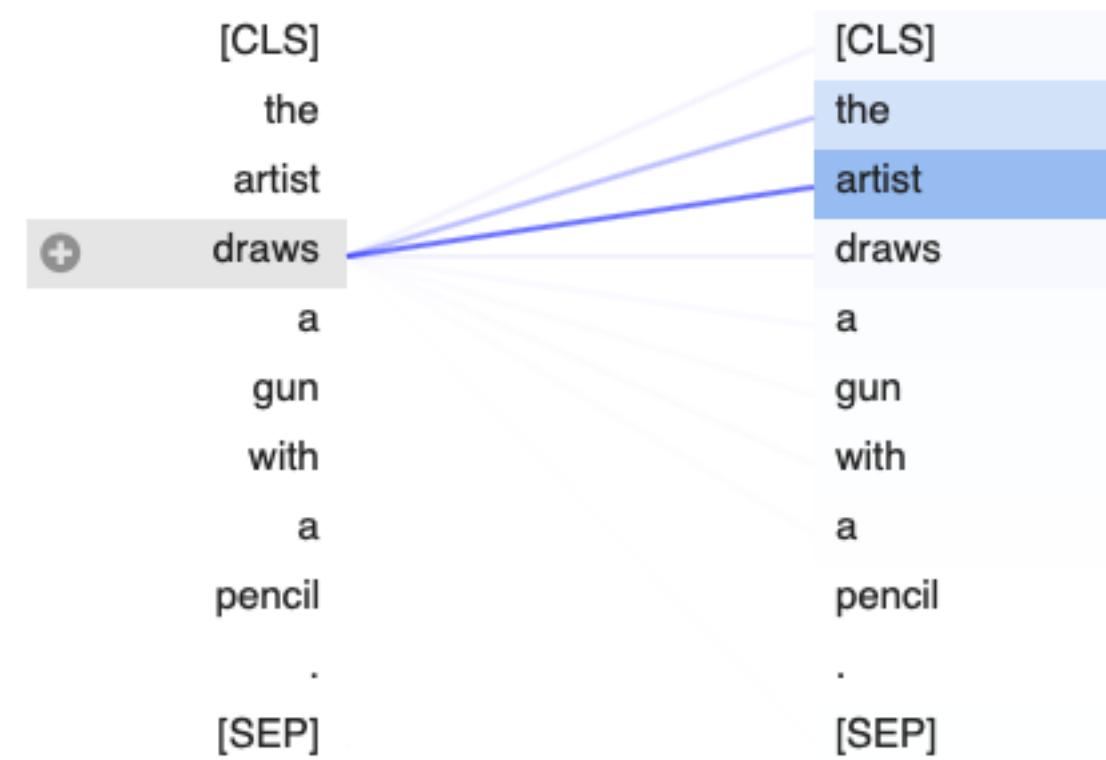
SKETCH

COMBINATION

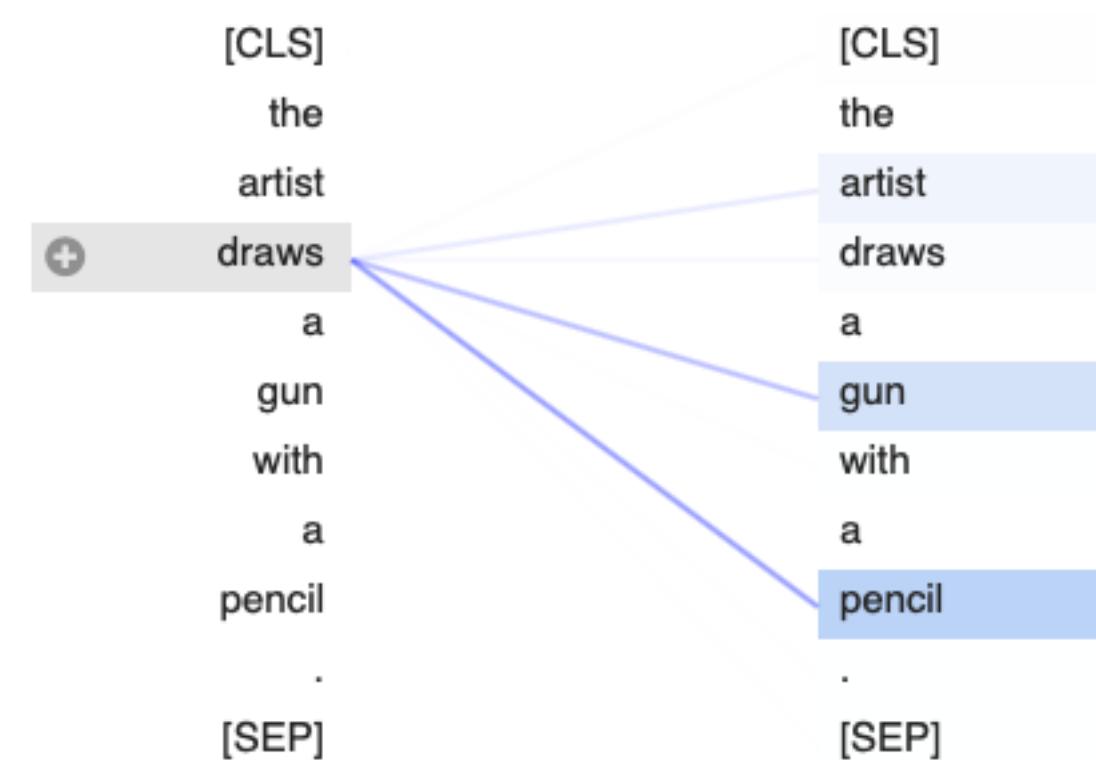
EACH HEAD QUERIES A DIFFERENT QUESTION
AND EXTRACTS A POSSIBLY DIFFERENT MEANING

SELF ATTENTION - ARTIST

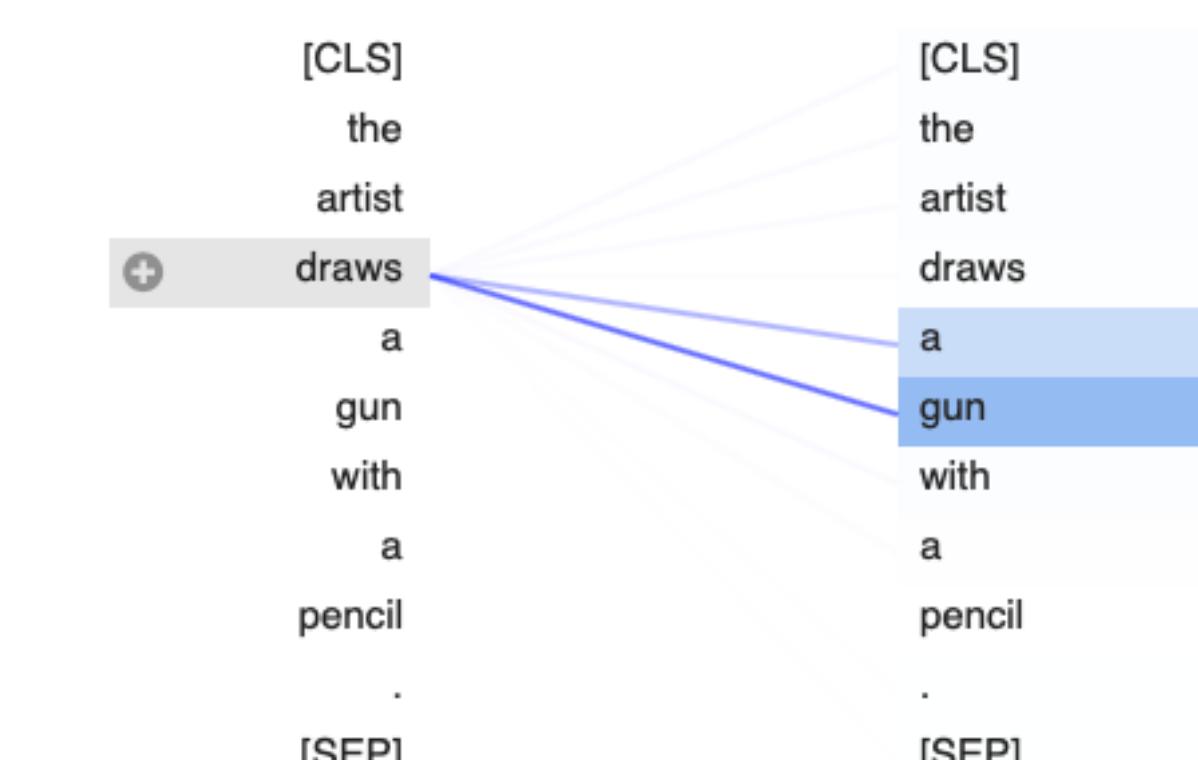
Layer: 0 Head: 3 Attention: Sentence A -> Sentence A



Layer: 0 Head: 5 Attention: Sentence A -> Sentence A



Layer: 0 Head: 10 Attention: Sentence A -> Sentence A



QUERY

WHO DRAWS?

KEYS

ARTIST

VALUES

SKETCH

COMBINATION / LINEAR LAYER

HOW IS IT DRAWN?

PENCIL

SKETCH

WHAT IS DRAWN?

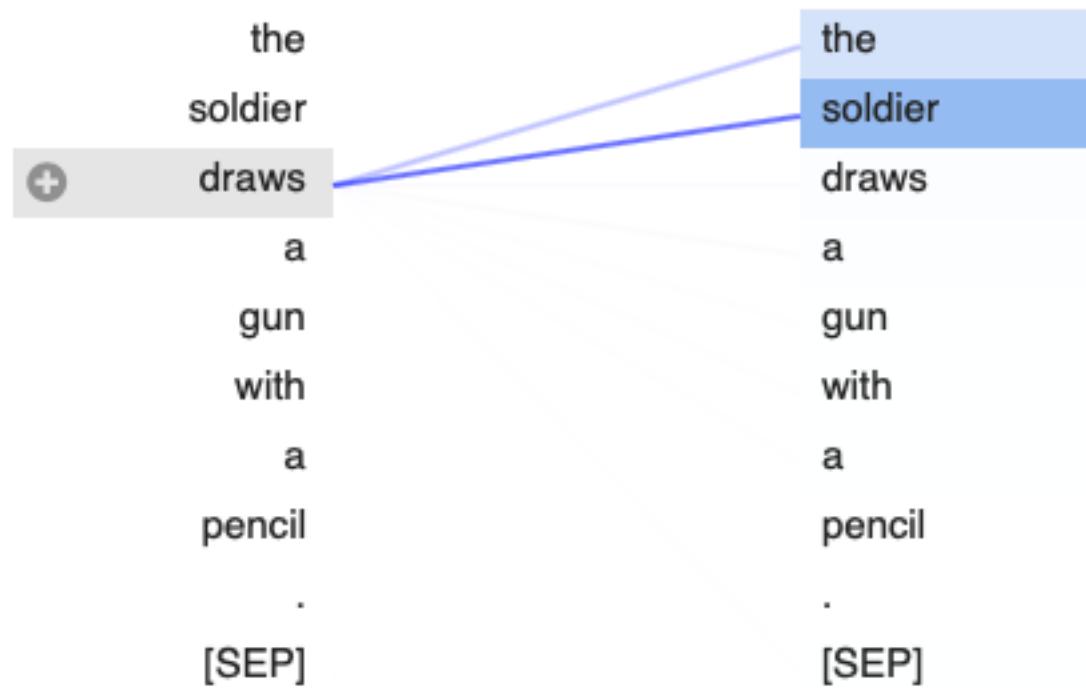
GUN

SHOOT

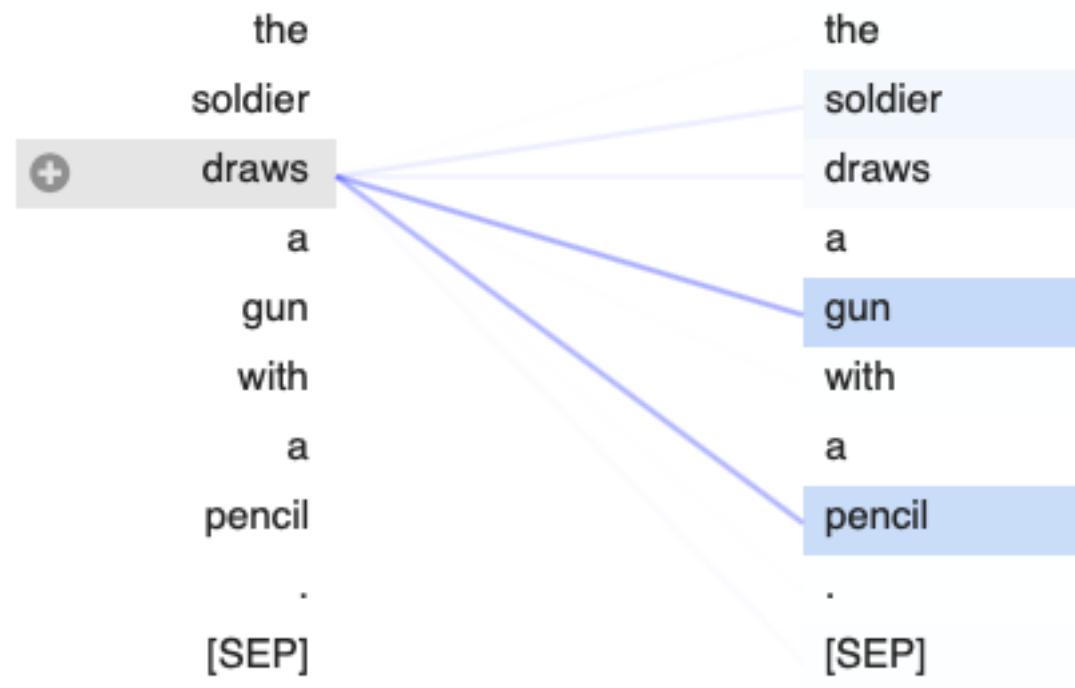
SKETCH

SELF ATTENTION - SOLDIER

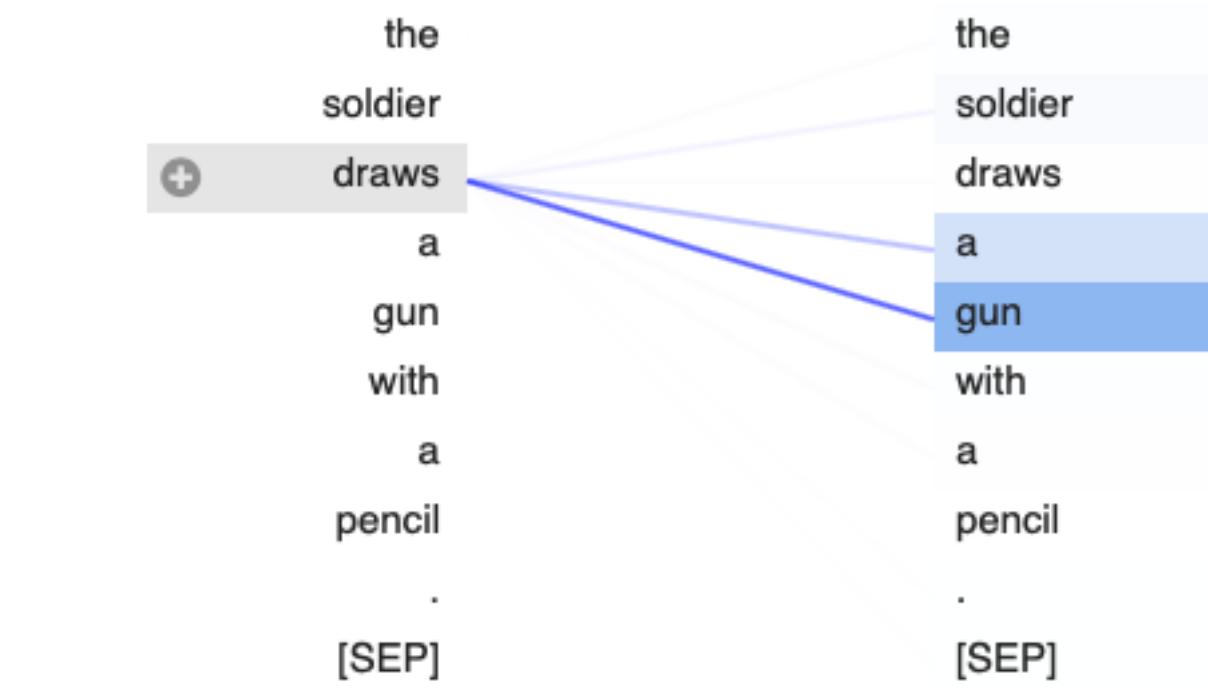
Layer: 0 Head: 3 Attention: Sentence B -> Sentence B



Layer: 0 Head: 5 Attention: Sentence B -> Sentence B



Layer: 0 Head: 10 Attention: Sentence B -> Sentence B



QUERY

WHO DRAWS?

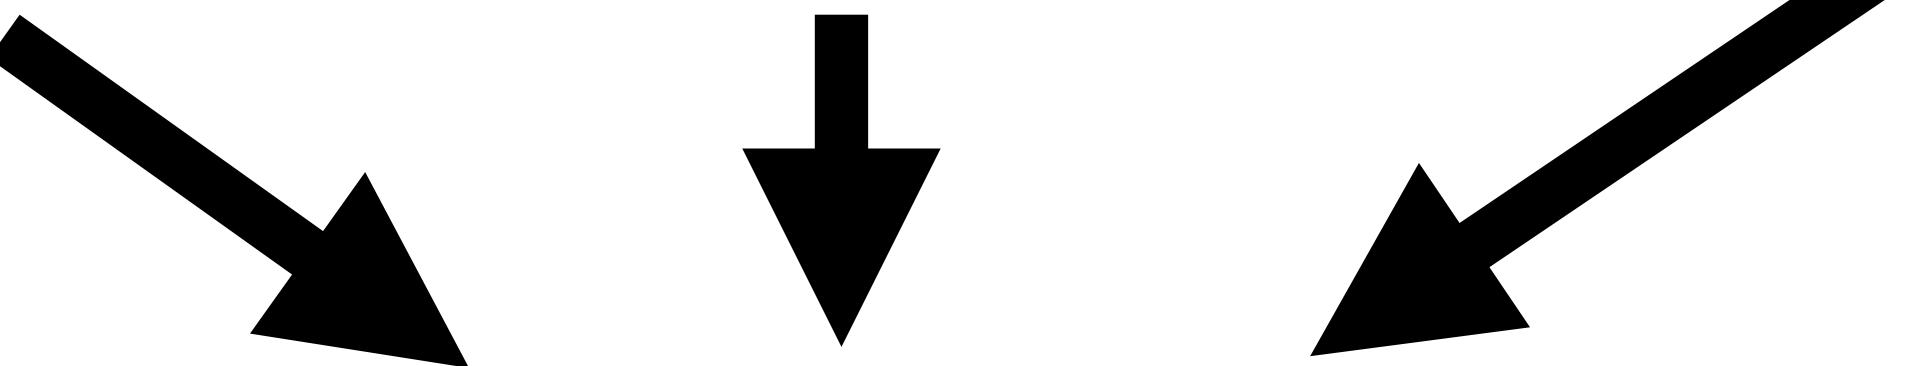
KEYS

SOLDIER

VALUES

SHOOT

COMBINATION / LINEAR LAYER



HOW IS IT DRAWN?

PENCIL

SKETCH

WHAT IS DRAWN?

GUN

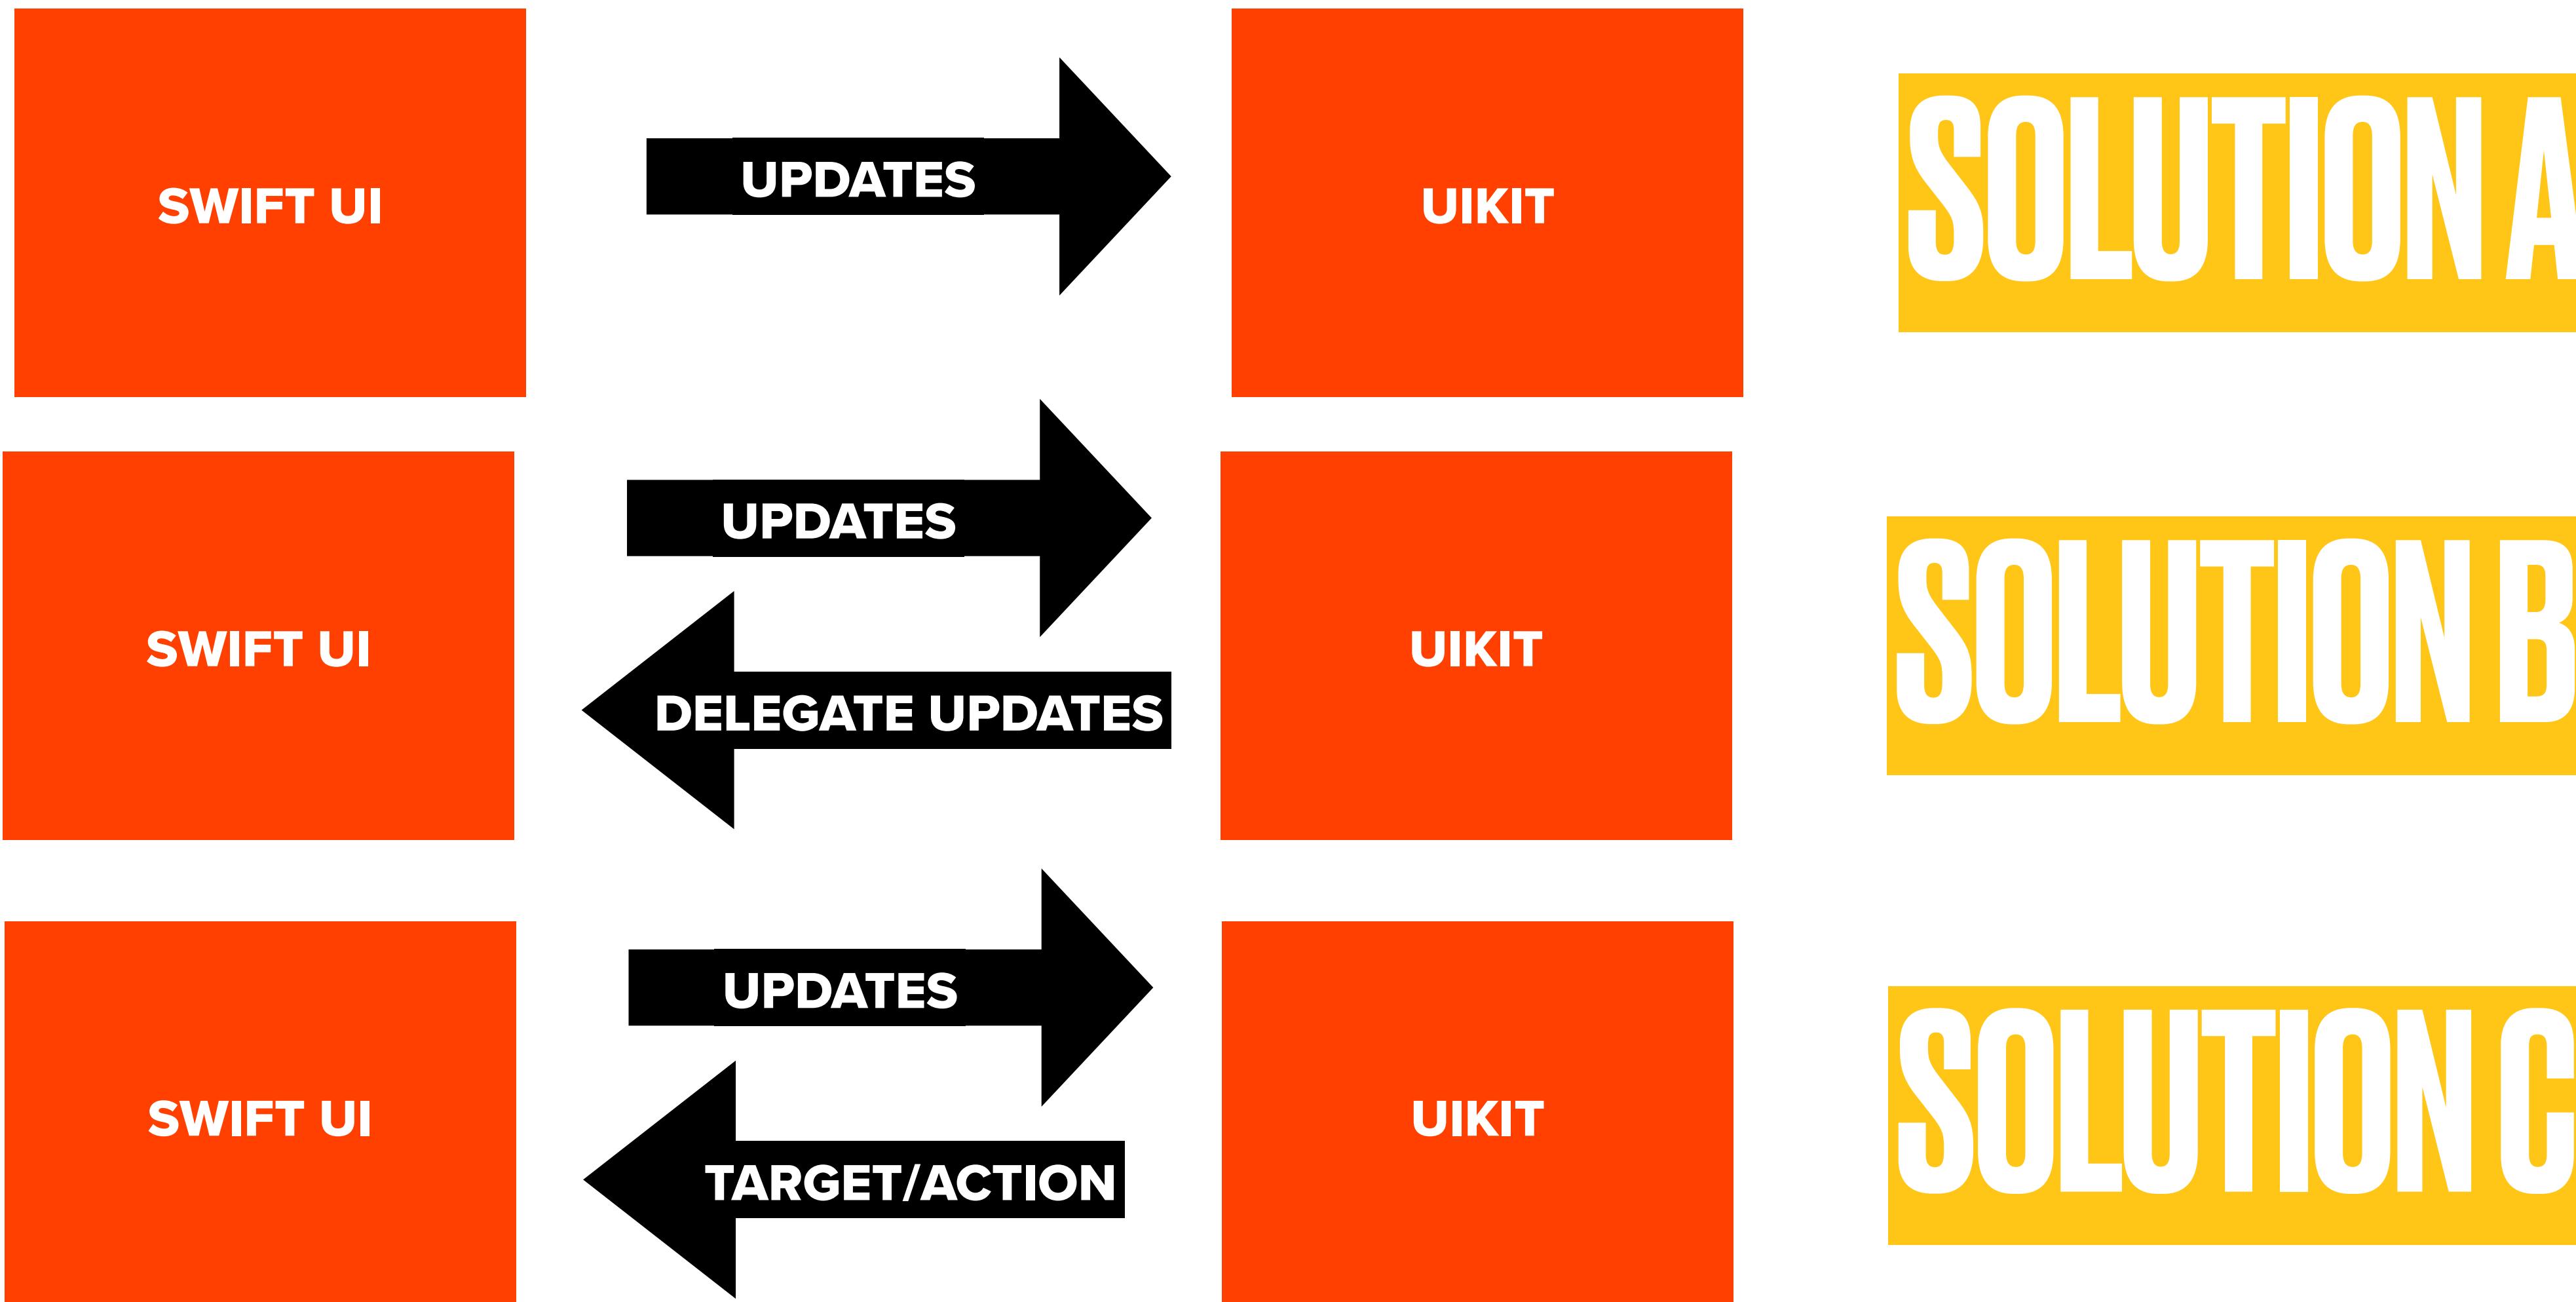
SHOOT

SHOOT

CODING 1: THE UI

CODING 2: THE PDF COMPONENT

HOW TO INTEGRATE A UIKIT COMPONENT IN SWIFTUI



SOLUTION A

- 1 Create a struct that holds the inputs as attributes with `@State` and `@Binding`
- 2 Extend the struct with the protocol
 - `UIView: UIViewRepresentable`
 - `UIViewController: UIViewControllerRepresentable`
- 3 Implement
 - `makeUIView`: return initialized UI View, use inputs for initializations
 - `updateUIView`: called when attributes change, so that you can react to those changes

SOLUTION B

- If you want to provide a delegate to your UIKit Component
- Implement delegate
 - Add @Binding attributes that you like to update
- Implement makeCoordinator
 - initialize delegate
 - pass @Binding attributes that you like to update
 - return delegate
- In makeUIViewController/makeUIView
 - register delegate with UIKit Component

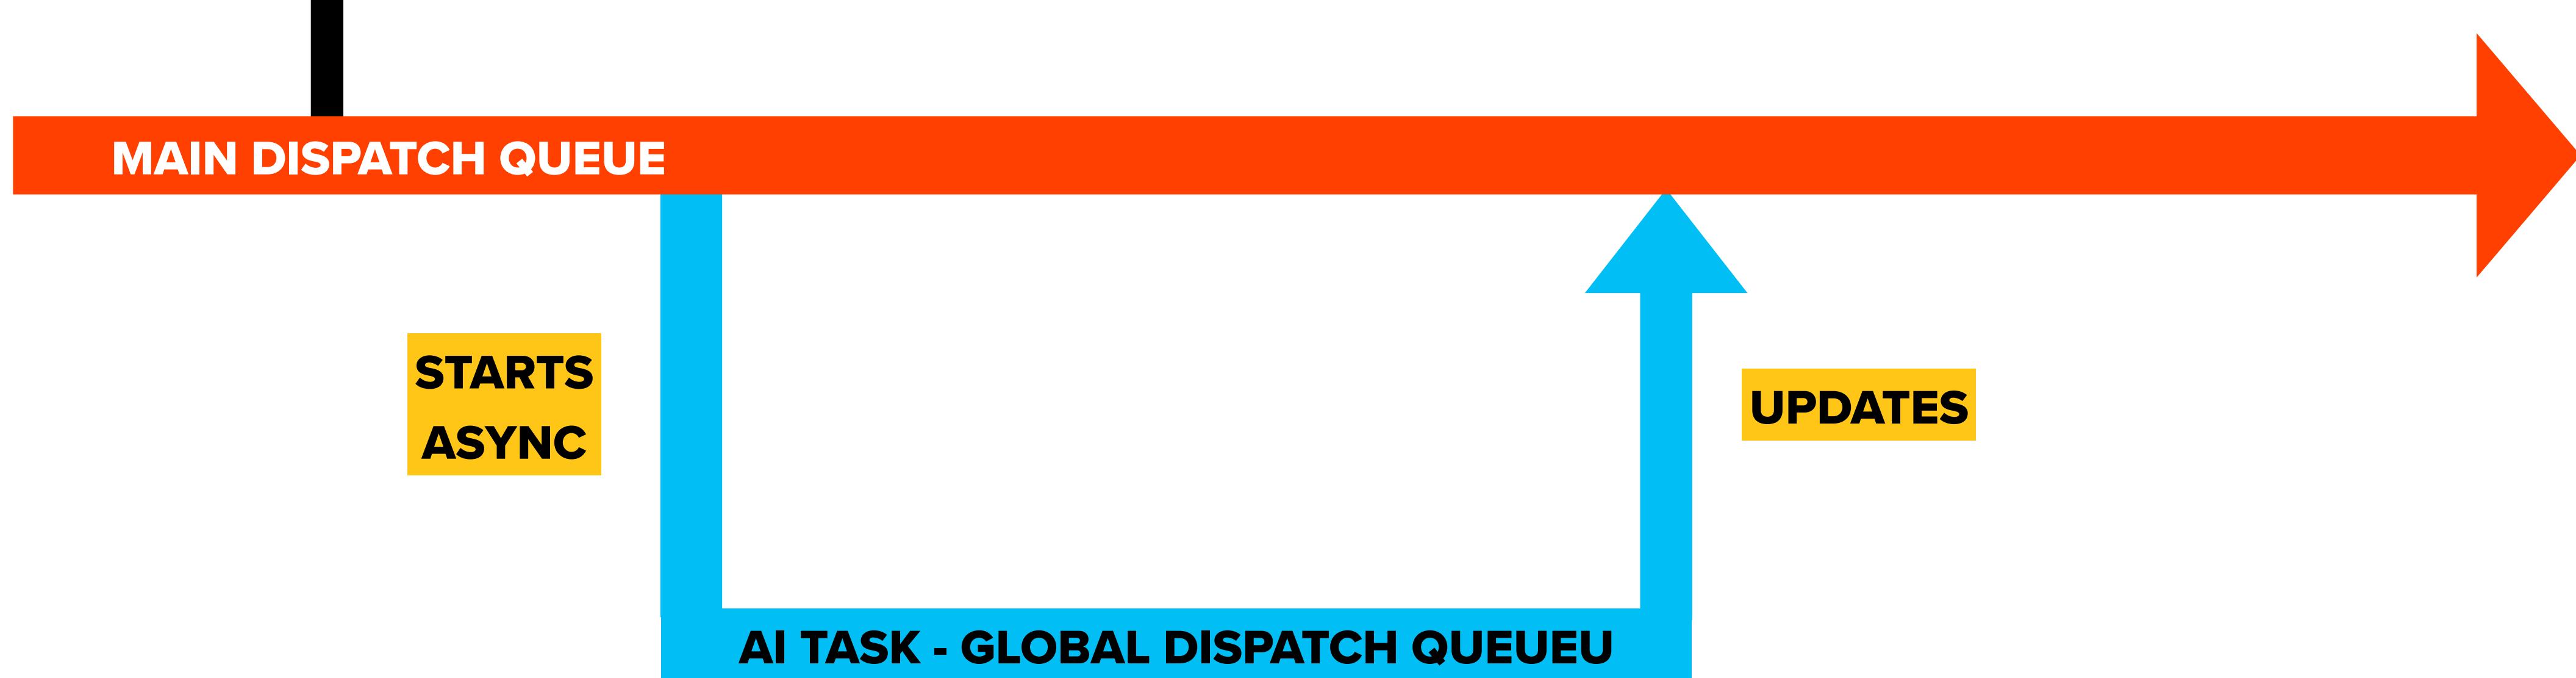
SOLUTION C

- If you need to update your Swift Component via Target Action Pattern
- Implement a new class inheriting from NSObject
 - In its init method, pass your UIViewRepresentable/UIViewControllerRepresentable
 - add a @objc function
 - with sender of UIKit Component type and
 - update UIViewRepresentable/UIViewControllerRepresentable based on UIKit Components changed value
- Implement makeCoordinator
 - initialize Coordinator with self
- In makeUIViewController/makeUIView
 - register your Coordinator @objc function for your update event such as .valueChanged

CODING 3: USING BERT WITH GCD

UNDERSTANDING THE GCD, JUST ENOUGH

USER INTERFACE - MAY NEVER FREEZE



QUALITY OF SERVICE - QOS - PRIORITY

USER-INTERACTIVE > USER-INITIATED > UTILITY > BACKGROUND

SINGLE AI TASK PATTERN

```
DispatchQueue.global(qos: .utility).async {
    // perform AI task
    DispatchQueue.main.async {
        //update UI
    }
}
```

UNDERSTANDING CONCURRENCY

- Global vs Private Dispatch Queues
 - Level of Concurrency: serial vs concurrent
 - Number of concurrent tasks
- Serial vs Concurrent Queues
 - Serial: one task at a time
 - Concurrent: multiple tasks at a time
- Sync vs Async
 - Sync: blocks the caller queue until task is finished (never on Main Queue)
 - Async: returns immediately, no blocking

DISPATCH GROUPS

- Run multiple AI tasks asynchronously
- Update UI when all of them are finished
- Use
 - DispatchGroup() to create group
 - for each task, group.enter() and group.leave()
 - make sure that you always leave the task by using defer
 - enter and leave serve as a concurrent counter for the group to know when all tasks are finished
 - use group.notify() to respond to the event that all tasks are done

```
we are async
finished: Text 2
finished: Text 3
finished: Text 1
all done
["Text 2 is the answer to Is the number of multiverses finite?", "Text 3 is the answer to Is the number of multiverses finite?", "Text 1 is the answer to Is the number of multiverses finite?"]
```

```
1 import UIKit
2
3 let aiTasks = DispatchGroup()
4 var results = [String]()
5
6 let inputs = ["Text 1", "Text 2", "Text 3" ]
7 let question = "Is the number of multiverses finite?"
8
9
10 for input in inputs {
11     aiTasks.enter()
12     DispatchQueue.global(qos: .utility).async {
13         defer { aiTasks.leave() }
14         // do your ai task
15         sleep(1)
16         print("finished: \(input)")
17         results.append("\(input) is the answer to \(question)")
18     }
19 }
20 print("we are async")
21 aiTasks.notify(queue: DispatchQueue.main) {
22     // update ui
23     print("all done")
24     print(results)
25 }
```

INTERMISSION

WHERE TO GO FROM HERE

THE COFFEE APP

COFFEE APP PURPOSE

- Pour Over Coffee - think of third wave coffee shops such as Blue Bottle
- At home you need to brew it with fairly precise measurements
- Each coffee may have a perfect sweet spot for each measure so you need to experiment a lot
- We will look at
 - Temperature
 - Grind size
- And build an app that functions as a diary for those measurements and star it based on coffee perfection



RESNET

PART 3: THE COFFEE APP

App Specifications

ResNet

Coding 1: The UI

Coding 2: Saving your Data - Designed for Privacy

Coding 3: Real-Time Camera Access

Coding 4: Integrating Vision Framework

Deep Residual Learning for Image Recognition

Kaiming He Xiangyu Zhang Shaoqing Ren Jian Sun
Microsoft Research
{kahe, v-xiangz, v-shren, jiansun}@microsoft.com

Deep residual learning for image recognition
K He, X Zhang, S Ren, J Sun - Proceedings of the IEEE ..., 2016 - openaccess.thecvf.com
Deeper neural networks are more difficult to train. We present a residual learning framework to ease the training of networks that are substantially deeper than those used previously. We explicitly reformulate the layers as learning residual functions with reference to the layer inputs, instead of learning unreferenced functions. We provide comprehensive empirical evidence showing that these residual networks are easier to optimize, and can gain accuracy from considerably increased depth. On the ImageNet dataset we evaluate residual ...



RESNET

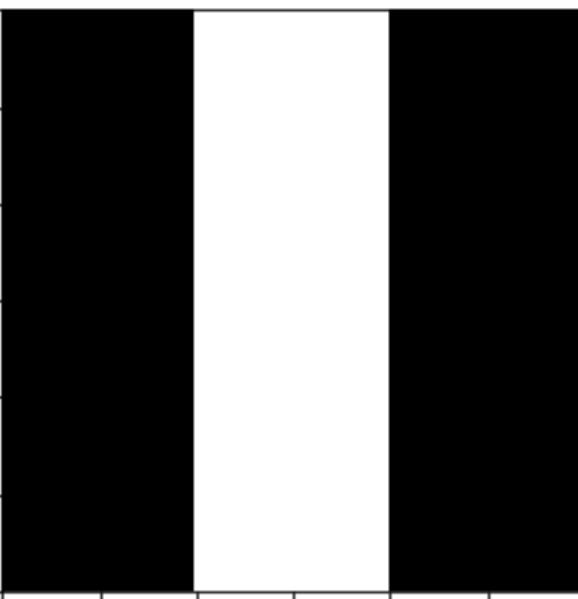
- **Background Knowledge:**
 - **CNN - Convolutional Neural Networks**
 - **CNNs learn filters with hierarchical learning**
 - **CNNs are Deep Neural Network**
 - **Deep Neural Networks are Neural Networks with many layers**
- **Problem:**
 - **Deep Neural Networks can suffer from Vanishing Gradient Problem**
 - **This causes the network not to learn any filters**
- **Solution:**
 - **ResNet addresses vanishing gradient problem with**
 - **Residual / Skip - Connections**

WHAT DOES A FILTER DO?



+

3X3 FILTER



CONVOLUTION →



INPUT

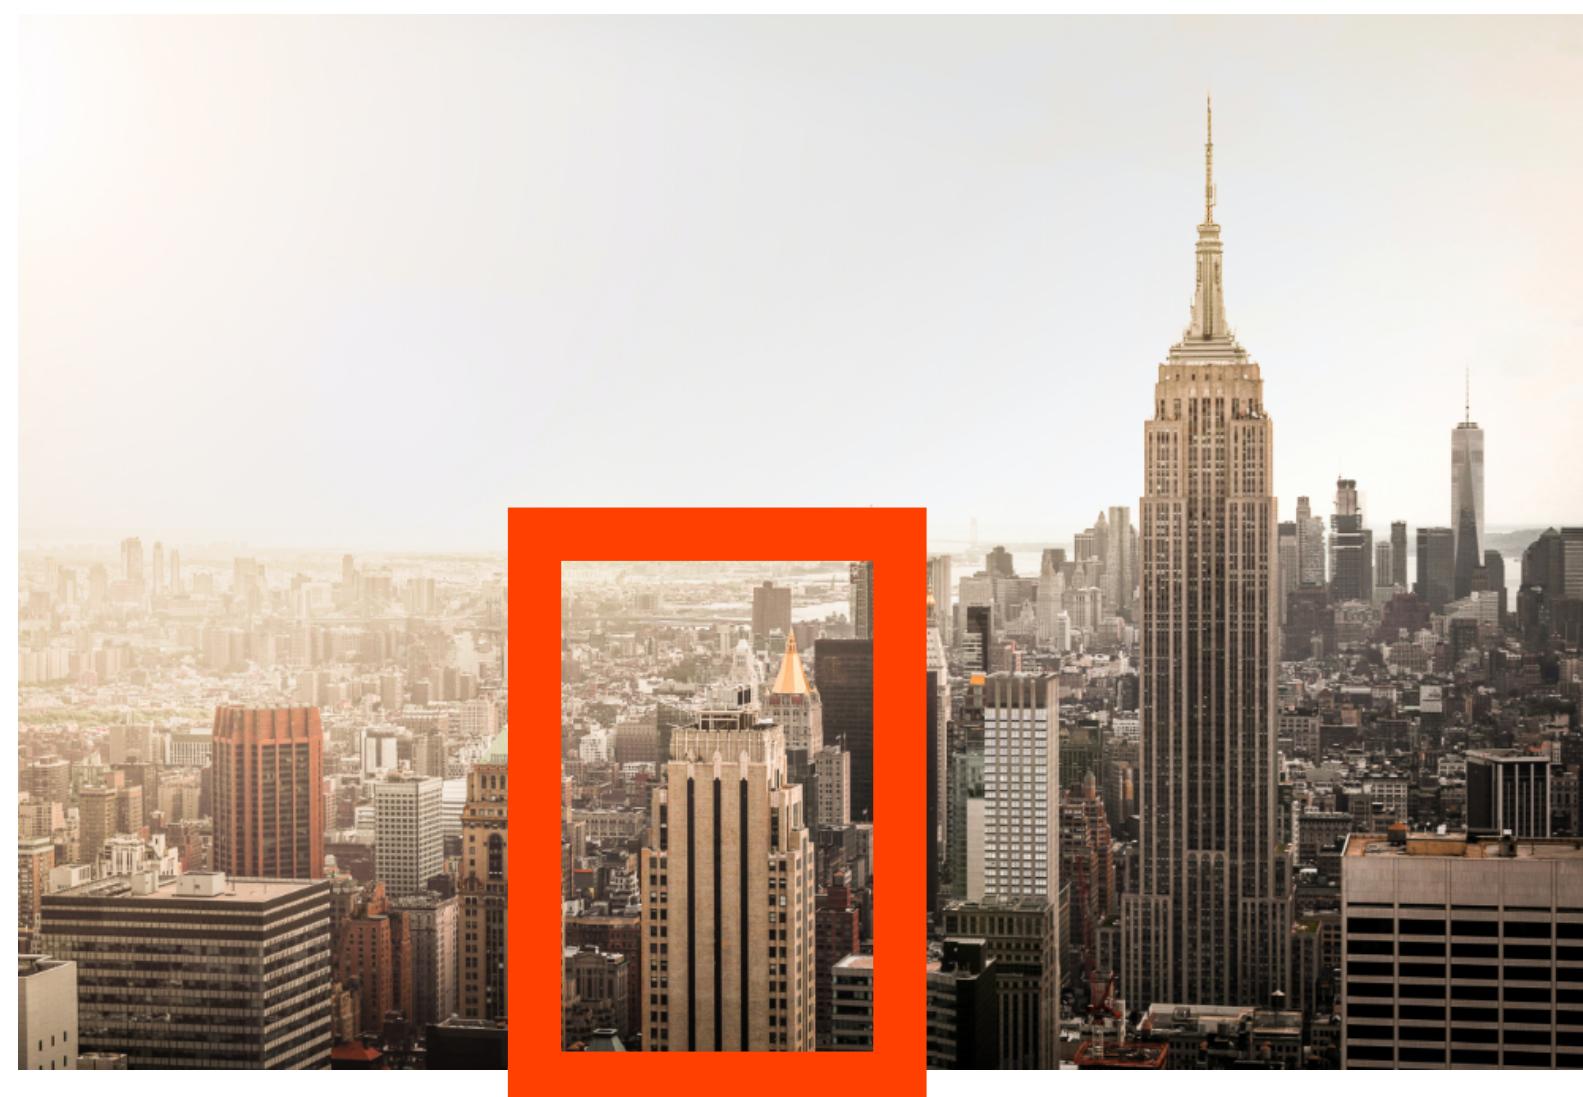
+

FILTER
FEATURE

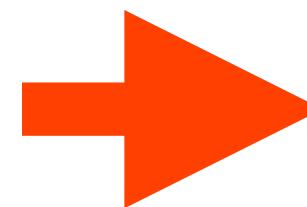
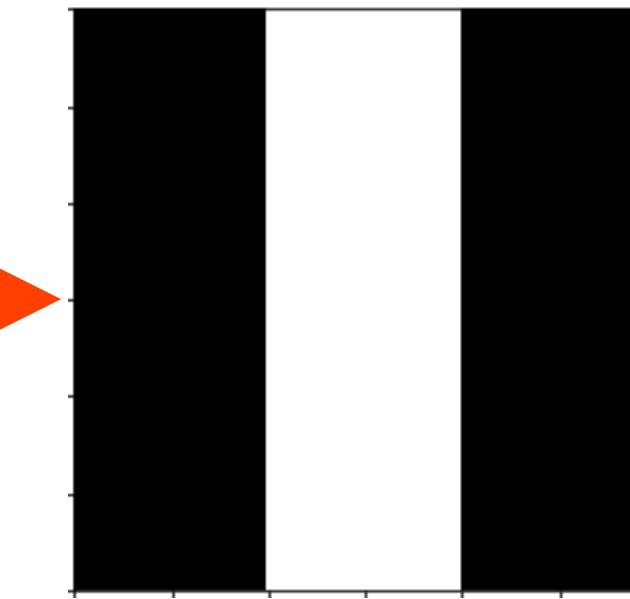
CONVOLUTION →

ACTIVATION MAP

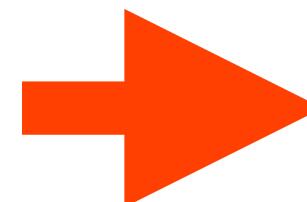
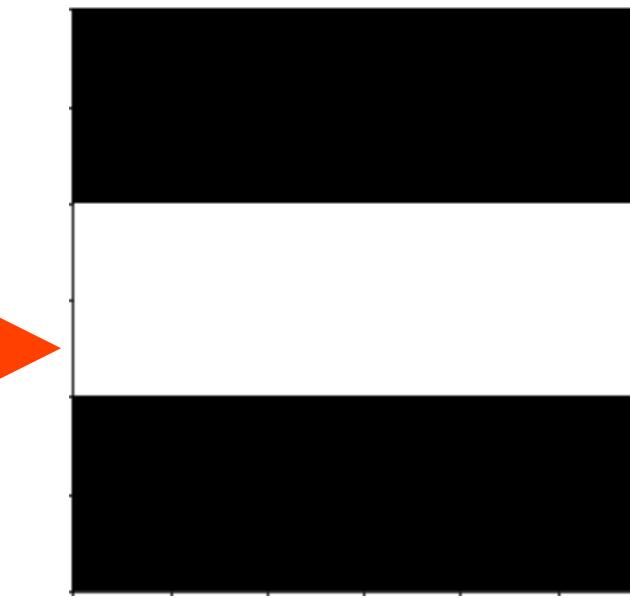
DIFFERENT FILTERS IN ACTION



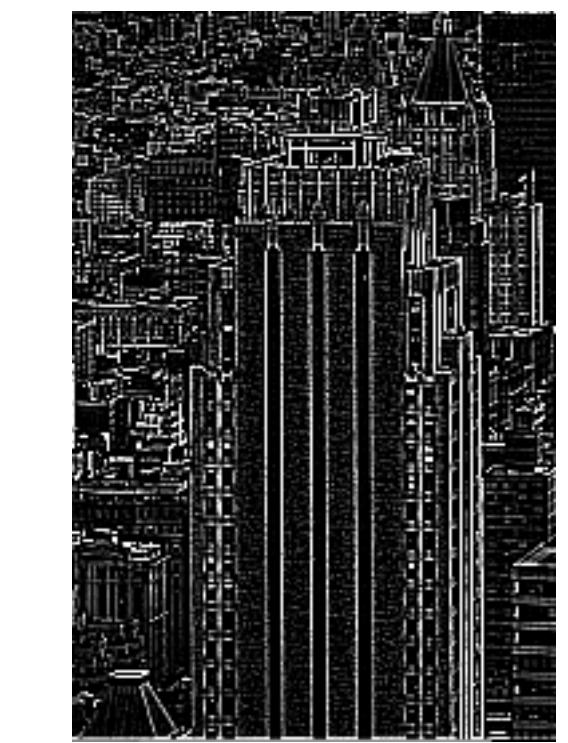
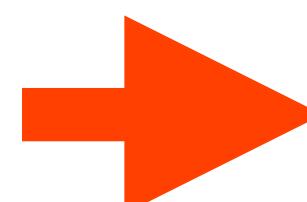
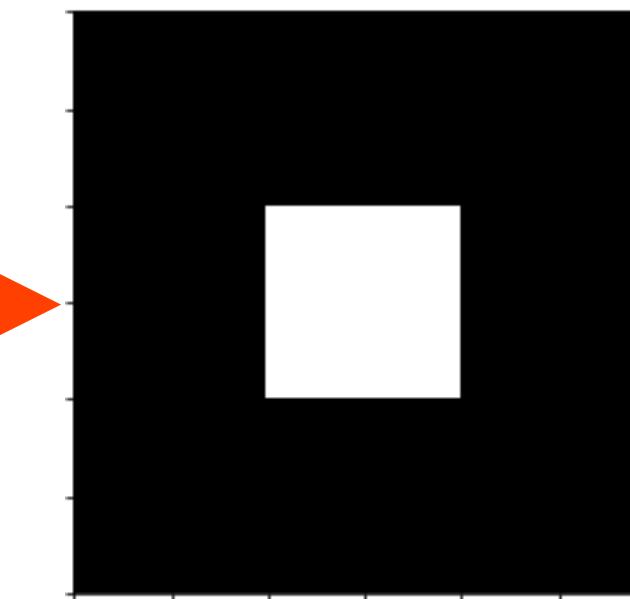
HORIZONTAL



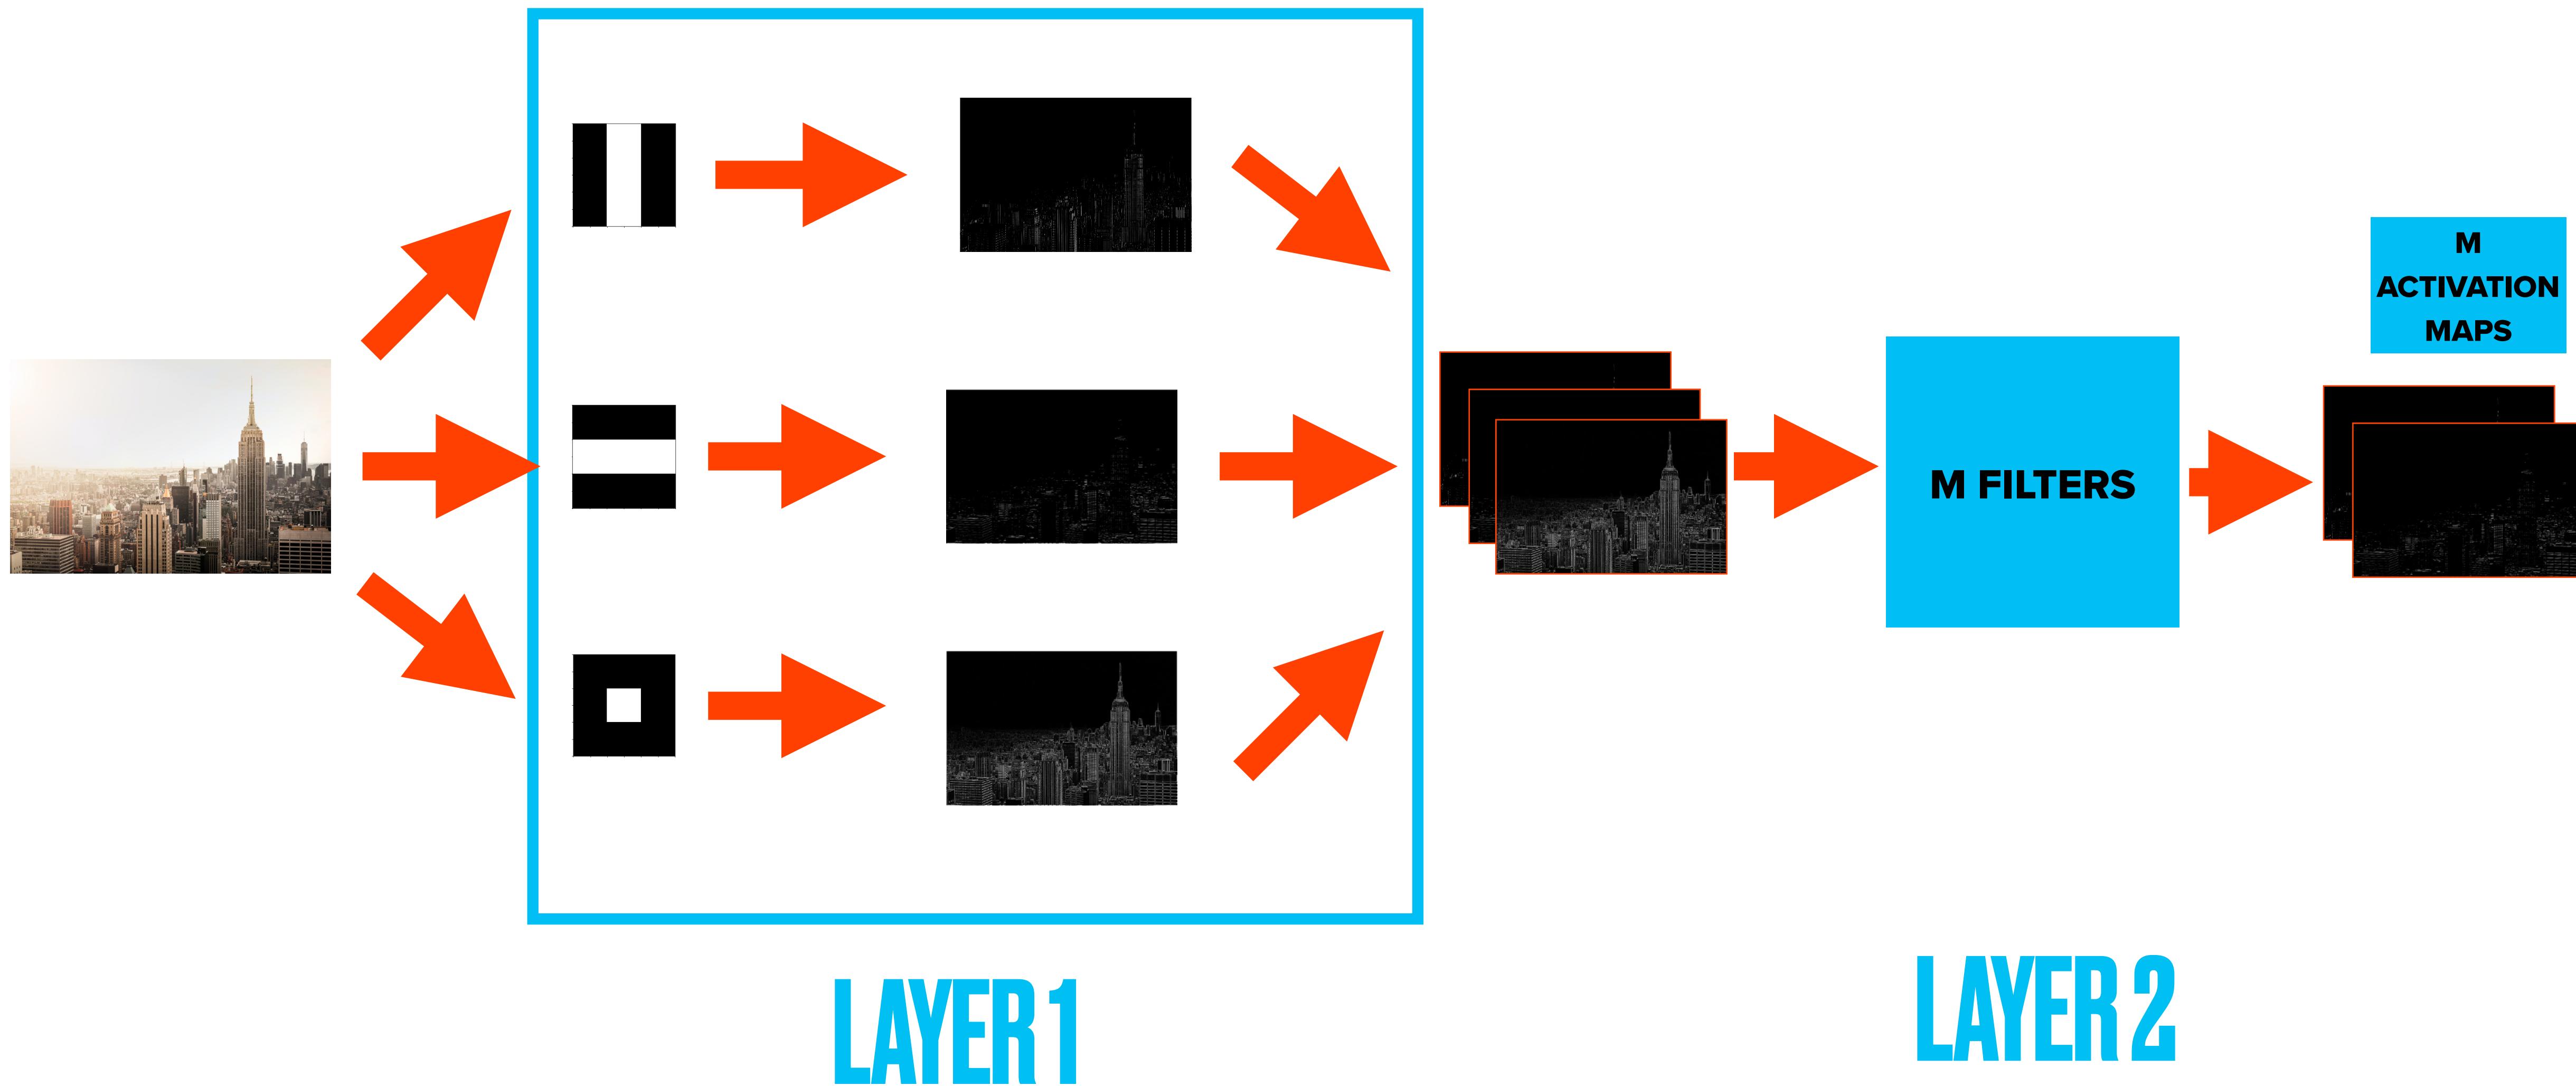
VERTICAL



EDGES



HIERARCHICAL LEARNING



HIERARCHICAL LEARNING - THE HOW

LITERAL → **ABSTRACT**

VERTICAL LINES
HORIZONTAL LINES
CURVES

PARALLEL LINES
CIRCLE
BOX

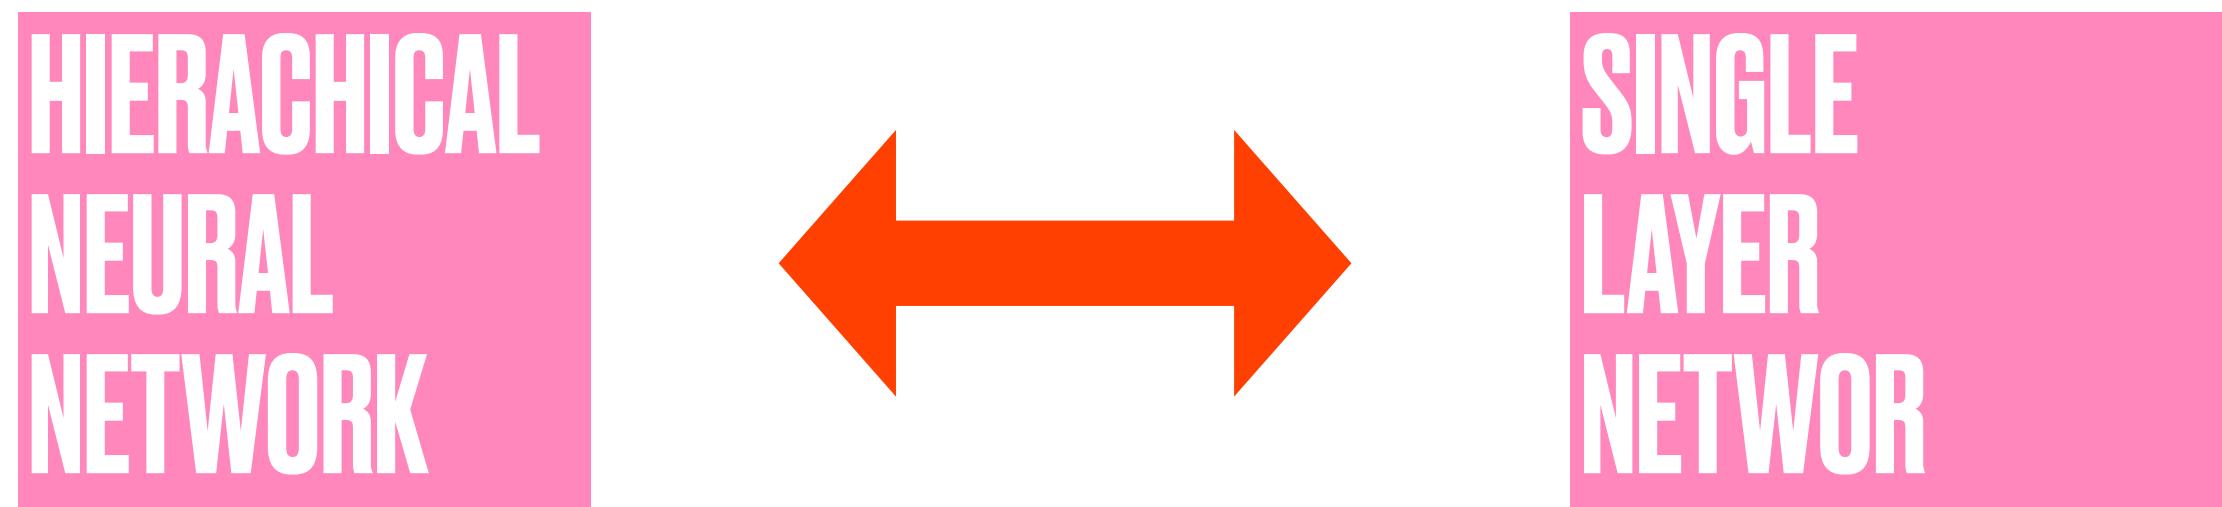
EYES
HANDS
STEERING WHEEL

HUMAN
CORVETTE
CAT

The more abstract the problem, the more layers we will need

HIERARCHICAL LEARNING - THE WHY

UNIVERSAL APPROXIMATION THEOREM

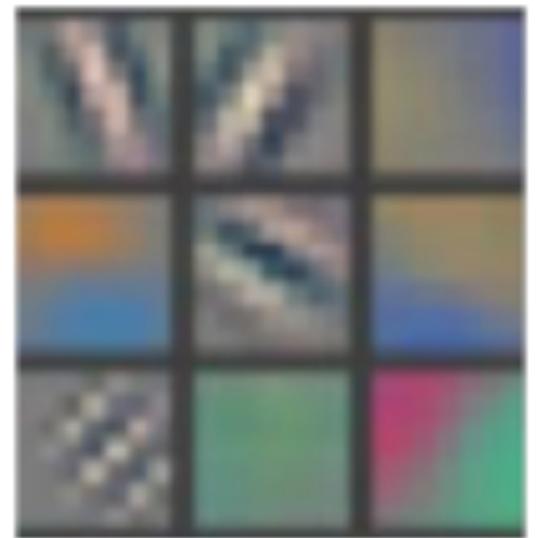


Reasons for Hierarchical Networks

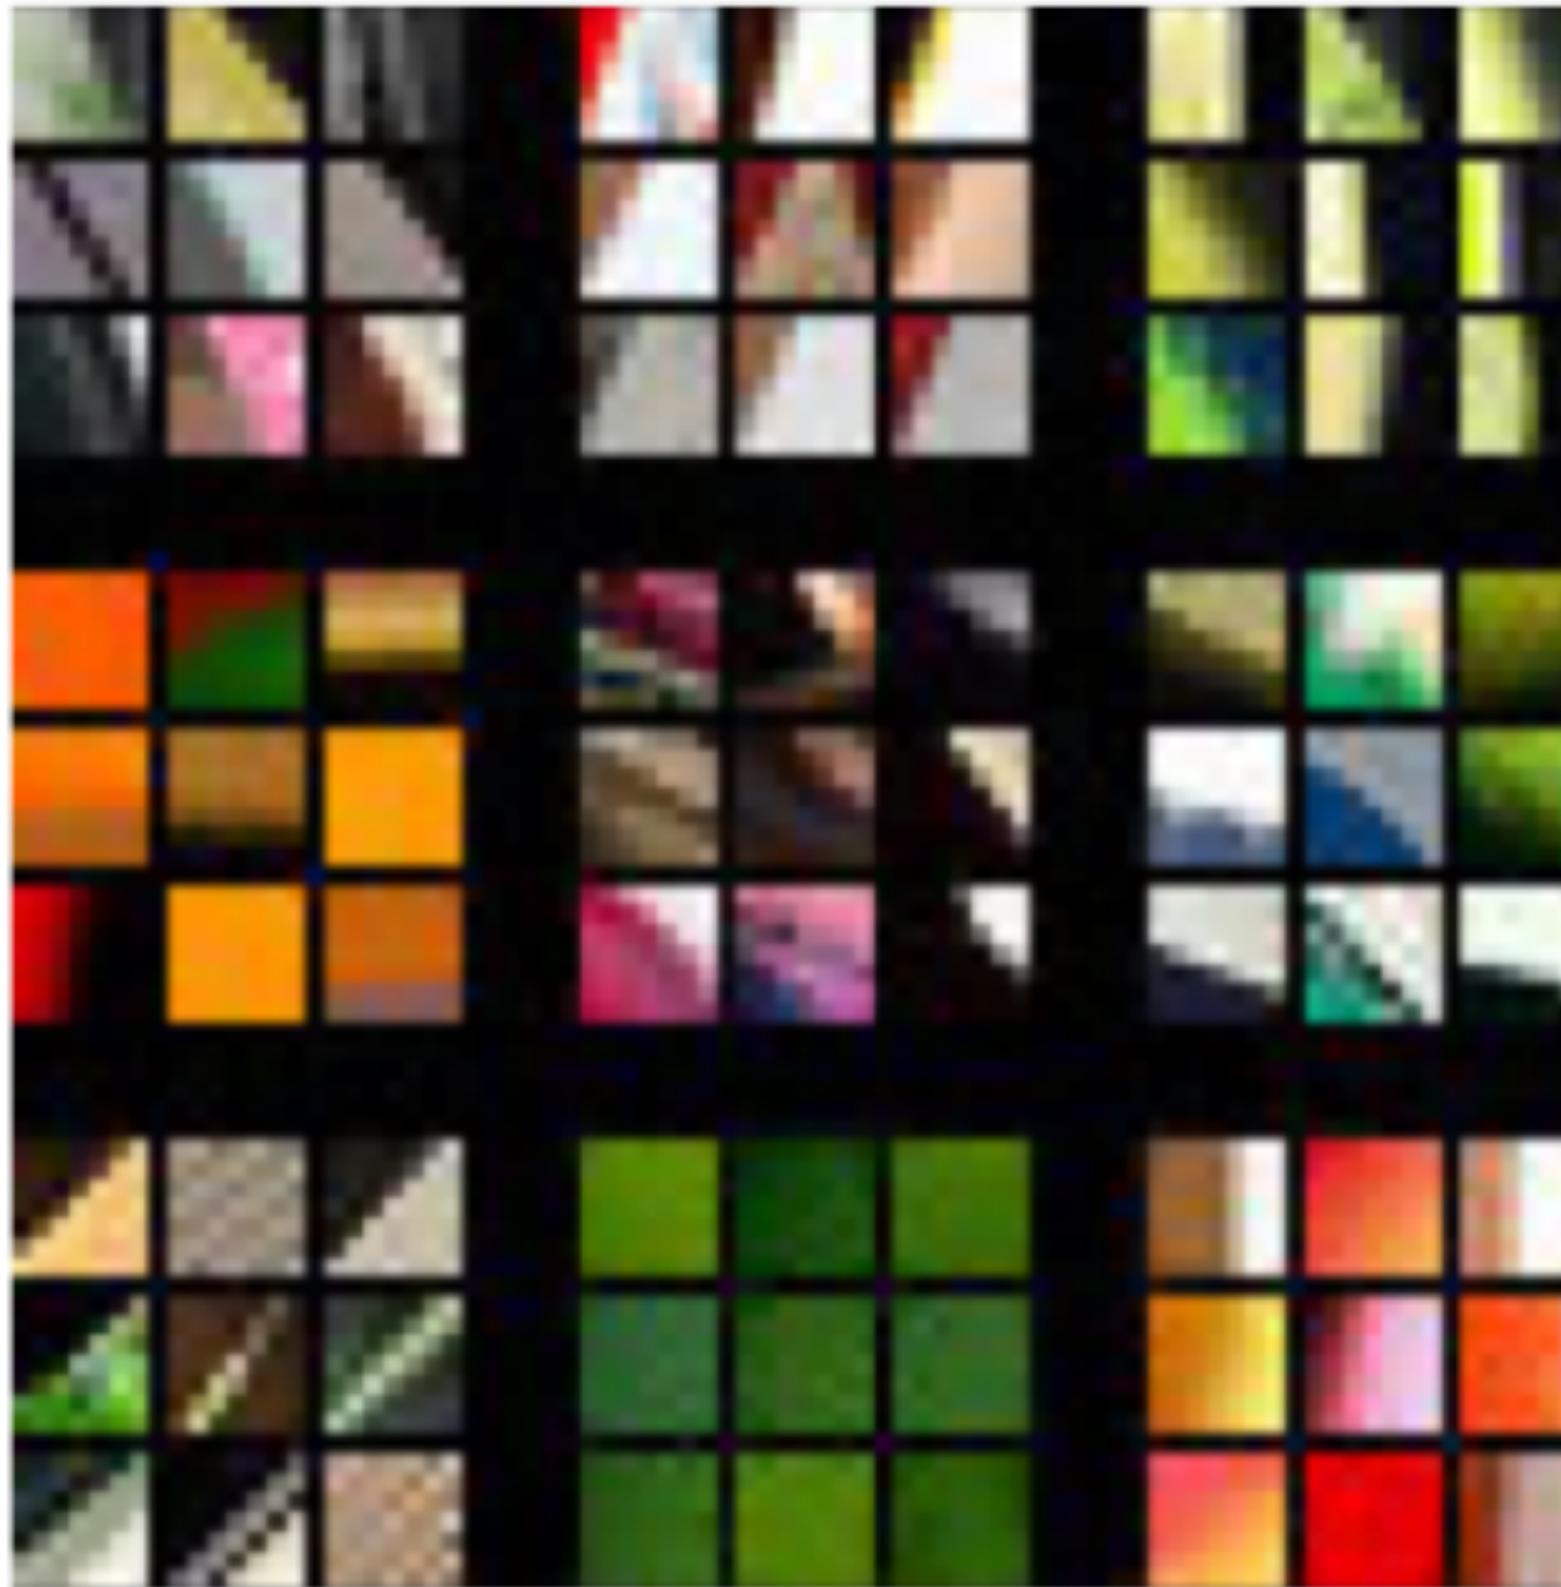
1. **Globally:** Fewer Neurons / Parameters
2. **Locally:** Fewer Neurons per Layer - better use of GPUs

**IN THE PAST WE CREATED FILTERS MANUALLY
CNNS LEARN FILTERS FROM DATA**

WHAT FILTERS DO CNNS LEARN?



Layer 1



Visualizing and Understanding Convolutional Networks

Matthew D. Zeiler

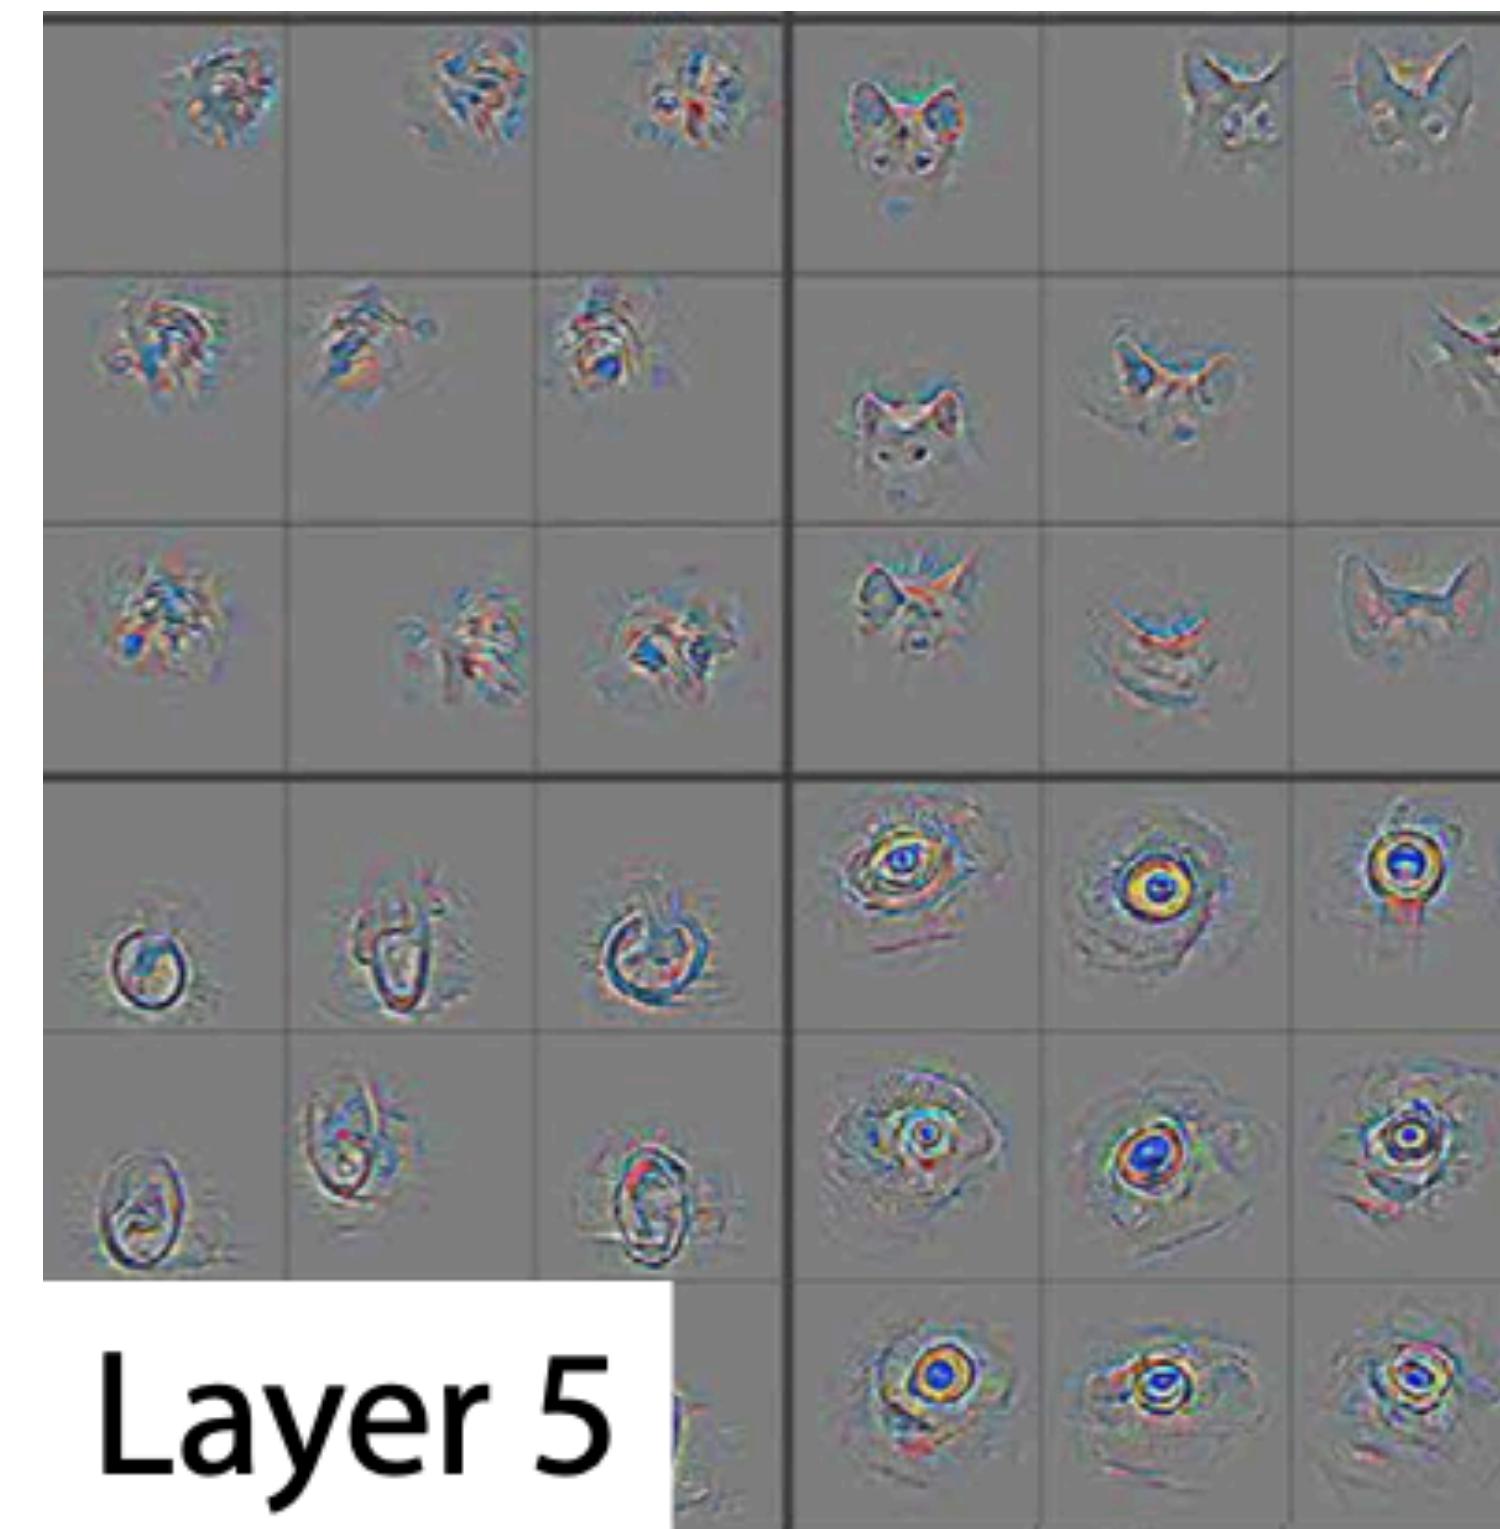
Dept. of Computer Science, Courant Institute, New York University

ZEILER@CS.NYU.EDU

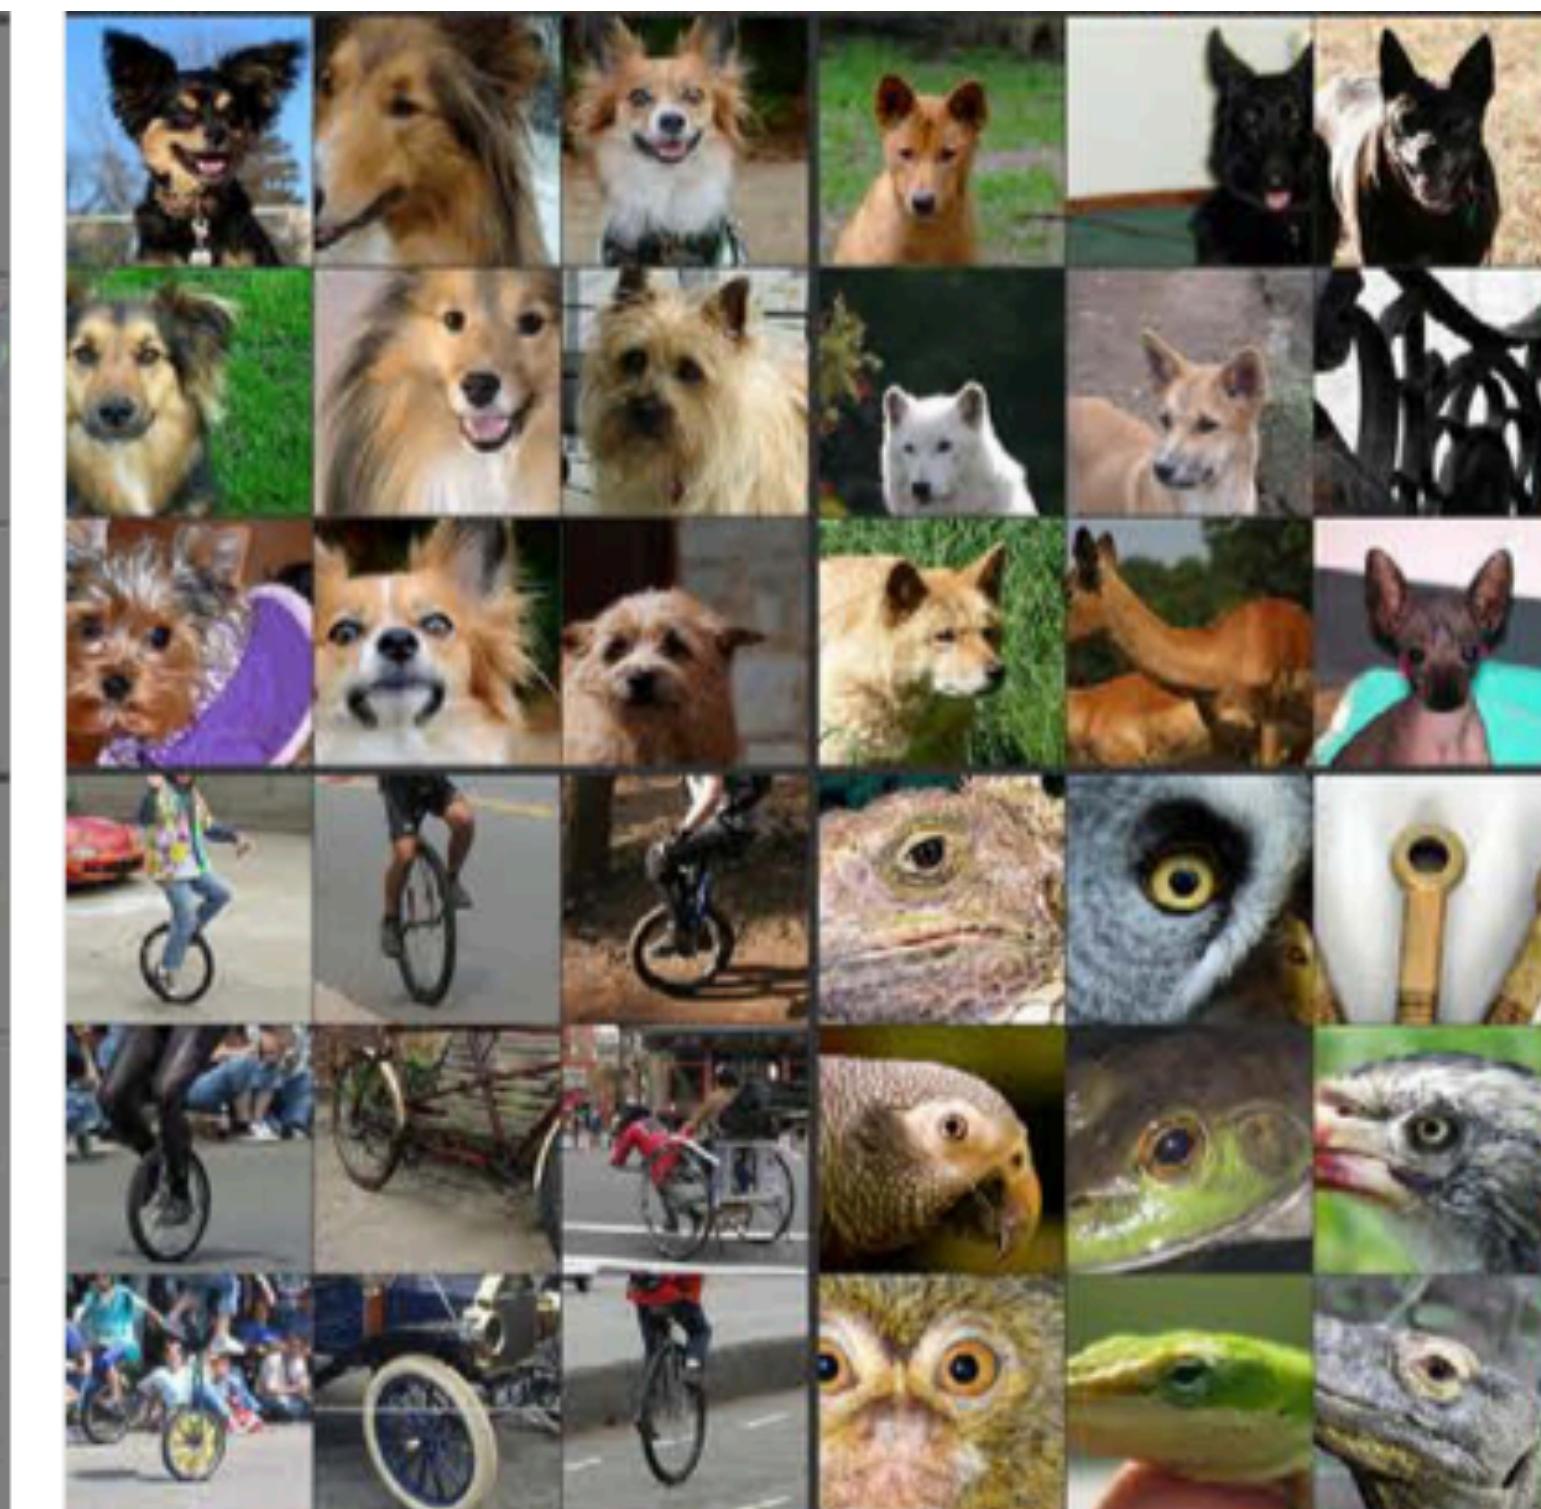
Rob Fergus

Dept. of Computer Science, Courant Institute, New York University

FERGUS@CS.NYU.EDU



Layer 5



VANISHING GRADIENT PROBLEM

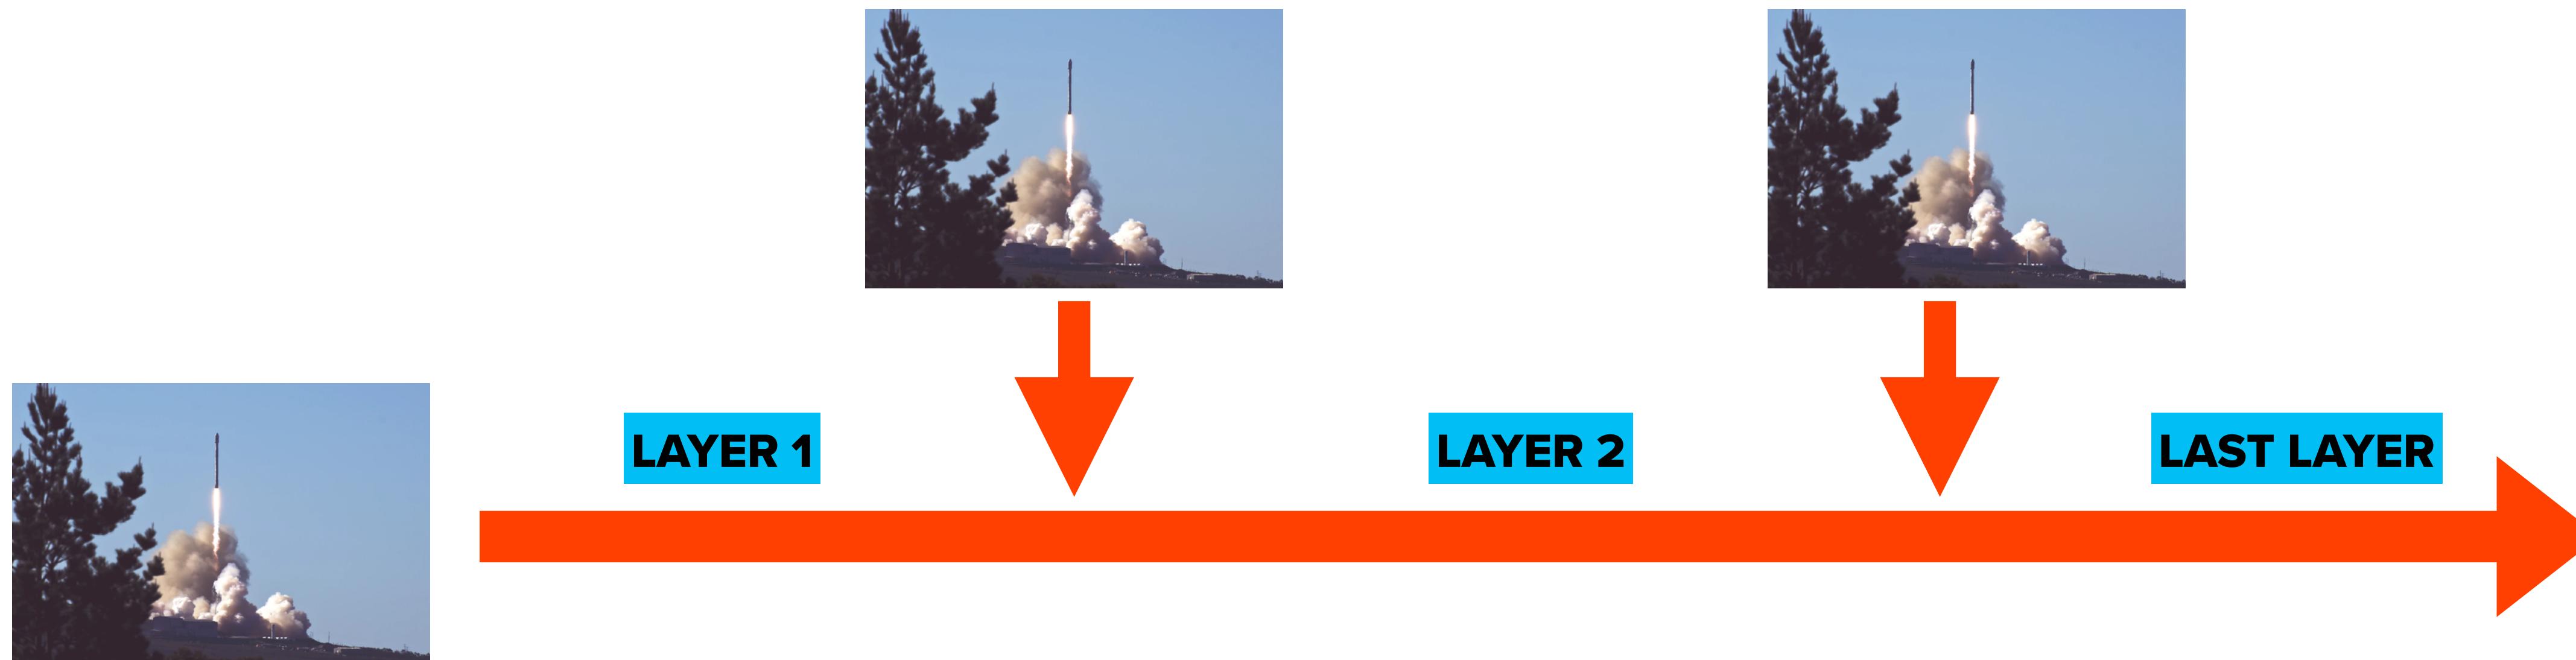
ALL INFORMATION → FILTERED → FEATURES

IMAGE
HAS CHROME
COLOR

CHROME
IS NOT A
FEATURE

IS IT A FALCON 9 ROCKET?

RESIDUAL / SKIP CONNECTIONS



1. Addresses the vanishing gradient problem due to proportional update of backpropogation
2. Transfer Learning: if you use a network not trained for your task, skip connection allow quicker convergence in practice because they can use the chrome just in the last layer and train only this layer anew instead of changing all previous layers

CODING 1: THE UI

DESIGN FOR “APPLE GLASSES”

CODING 2: SAVING YOUR DATA DESIGNED FOR PRIVACY

WHERE TO STORE USER IMAGES

■ iOS Data Storage Guidelines

Storing Your App's Data Efficiently

To ensure that backups are as efficient as possible, store your app's data according to the following guidelines:

1. Only documents and other data that is user-generated, or that cannot otherwise be recreated by your application, should be stored in the <Application_Home>/Documents directory and will be automatically backed up by iCloud.
2. Data that can be downloaded again or regenerated should be stored in the

■ The images are user-generated

■ Saving them in the documents directory allows backup in iCloud

```
FileManager.default.urls(for: .documentDirectory, in: .userDomainMask).first
```

```
file:///Users/stefan/Library/Developer/XCPGDevices/88E
```

JSON SERIALIZATION

- 1: Adopt the Codable protocol on the struct/class you want to save or load from disk
- 2: Do you have any custom types as attributes?
 - Then those custom types also need to adopt the Codable protocol
 - Native types already support it
- To save to disk:

```
let encoder = JSONEncoder()
do {
    let beansData = try encoder.encode(coffeeBeans)
    try beansData.write(to: collectionUrl, options: .atomicWrite)
} catch let error {
    // here we need to add proper error handling in production applications
    print(error)
}
```

- To load from disk:

```
let decoder = JSONDecoder()

do {
    let beansData = try Data(contentsOf: collectionUrl)
    coffeeBeans = try decoder.decode([CoffeeBean].self, from: beansData)
} catch let error {
    // here we need to add proper error handling in production applications
    print(error)
}
```

CODING 3: REAL-TIME CAMERA ACCESS

CAMERA RECORDING SESSION / CAMERA CONTROLLER

- Inherit from **AVCaptureVideoDataOutputSampleBufferDelegate**
- Setup an **AVCaptureSession**, add an input device and configure them as needed
- Queue **AVCaptureVideoDataOutput** to your **VNImageRequestHandler**

```
let session = AVCaptureSession()
private let videoDataOutput = AVCaptureVideoDataOutput()
let videoDataOutputQueue = DispatchQueue(label: "VideoDataOutput", qos: .userInitiated, attributes: [],
autoreleaseFrequency: .workItem)
let videoDevice = AVCaptureDevice.DiscoverySession(deviceTypes: [.builtInWideAngleCamera],
mediaType: .video, position: .back).devices.first

videoDataOutput.setSampleBufferDelegate(self, queue: videoDataOutputQueue)

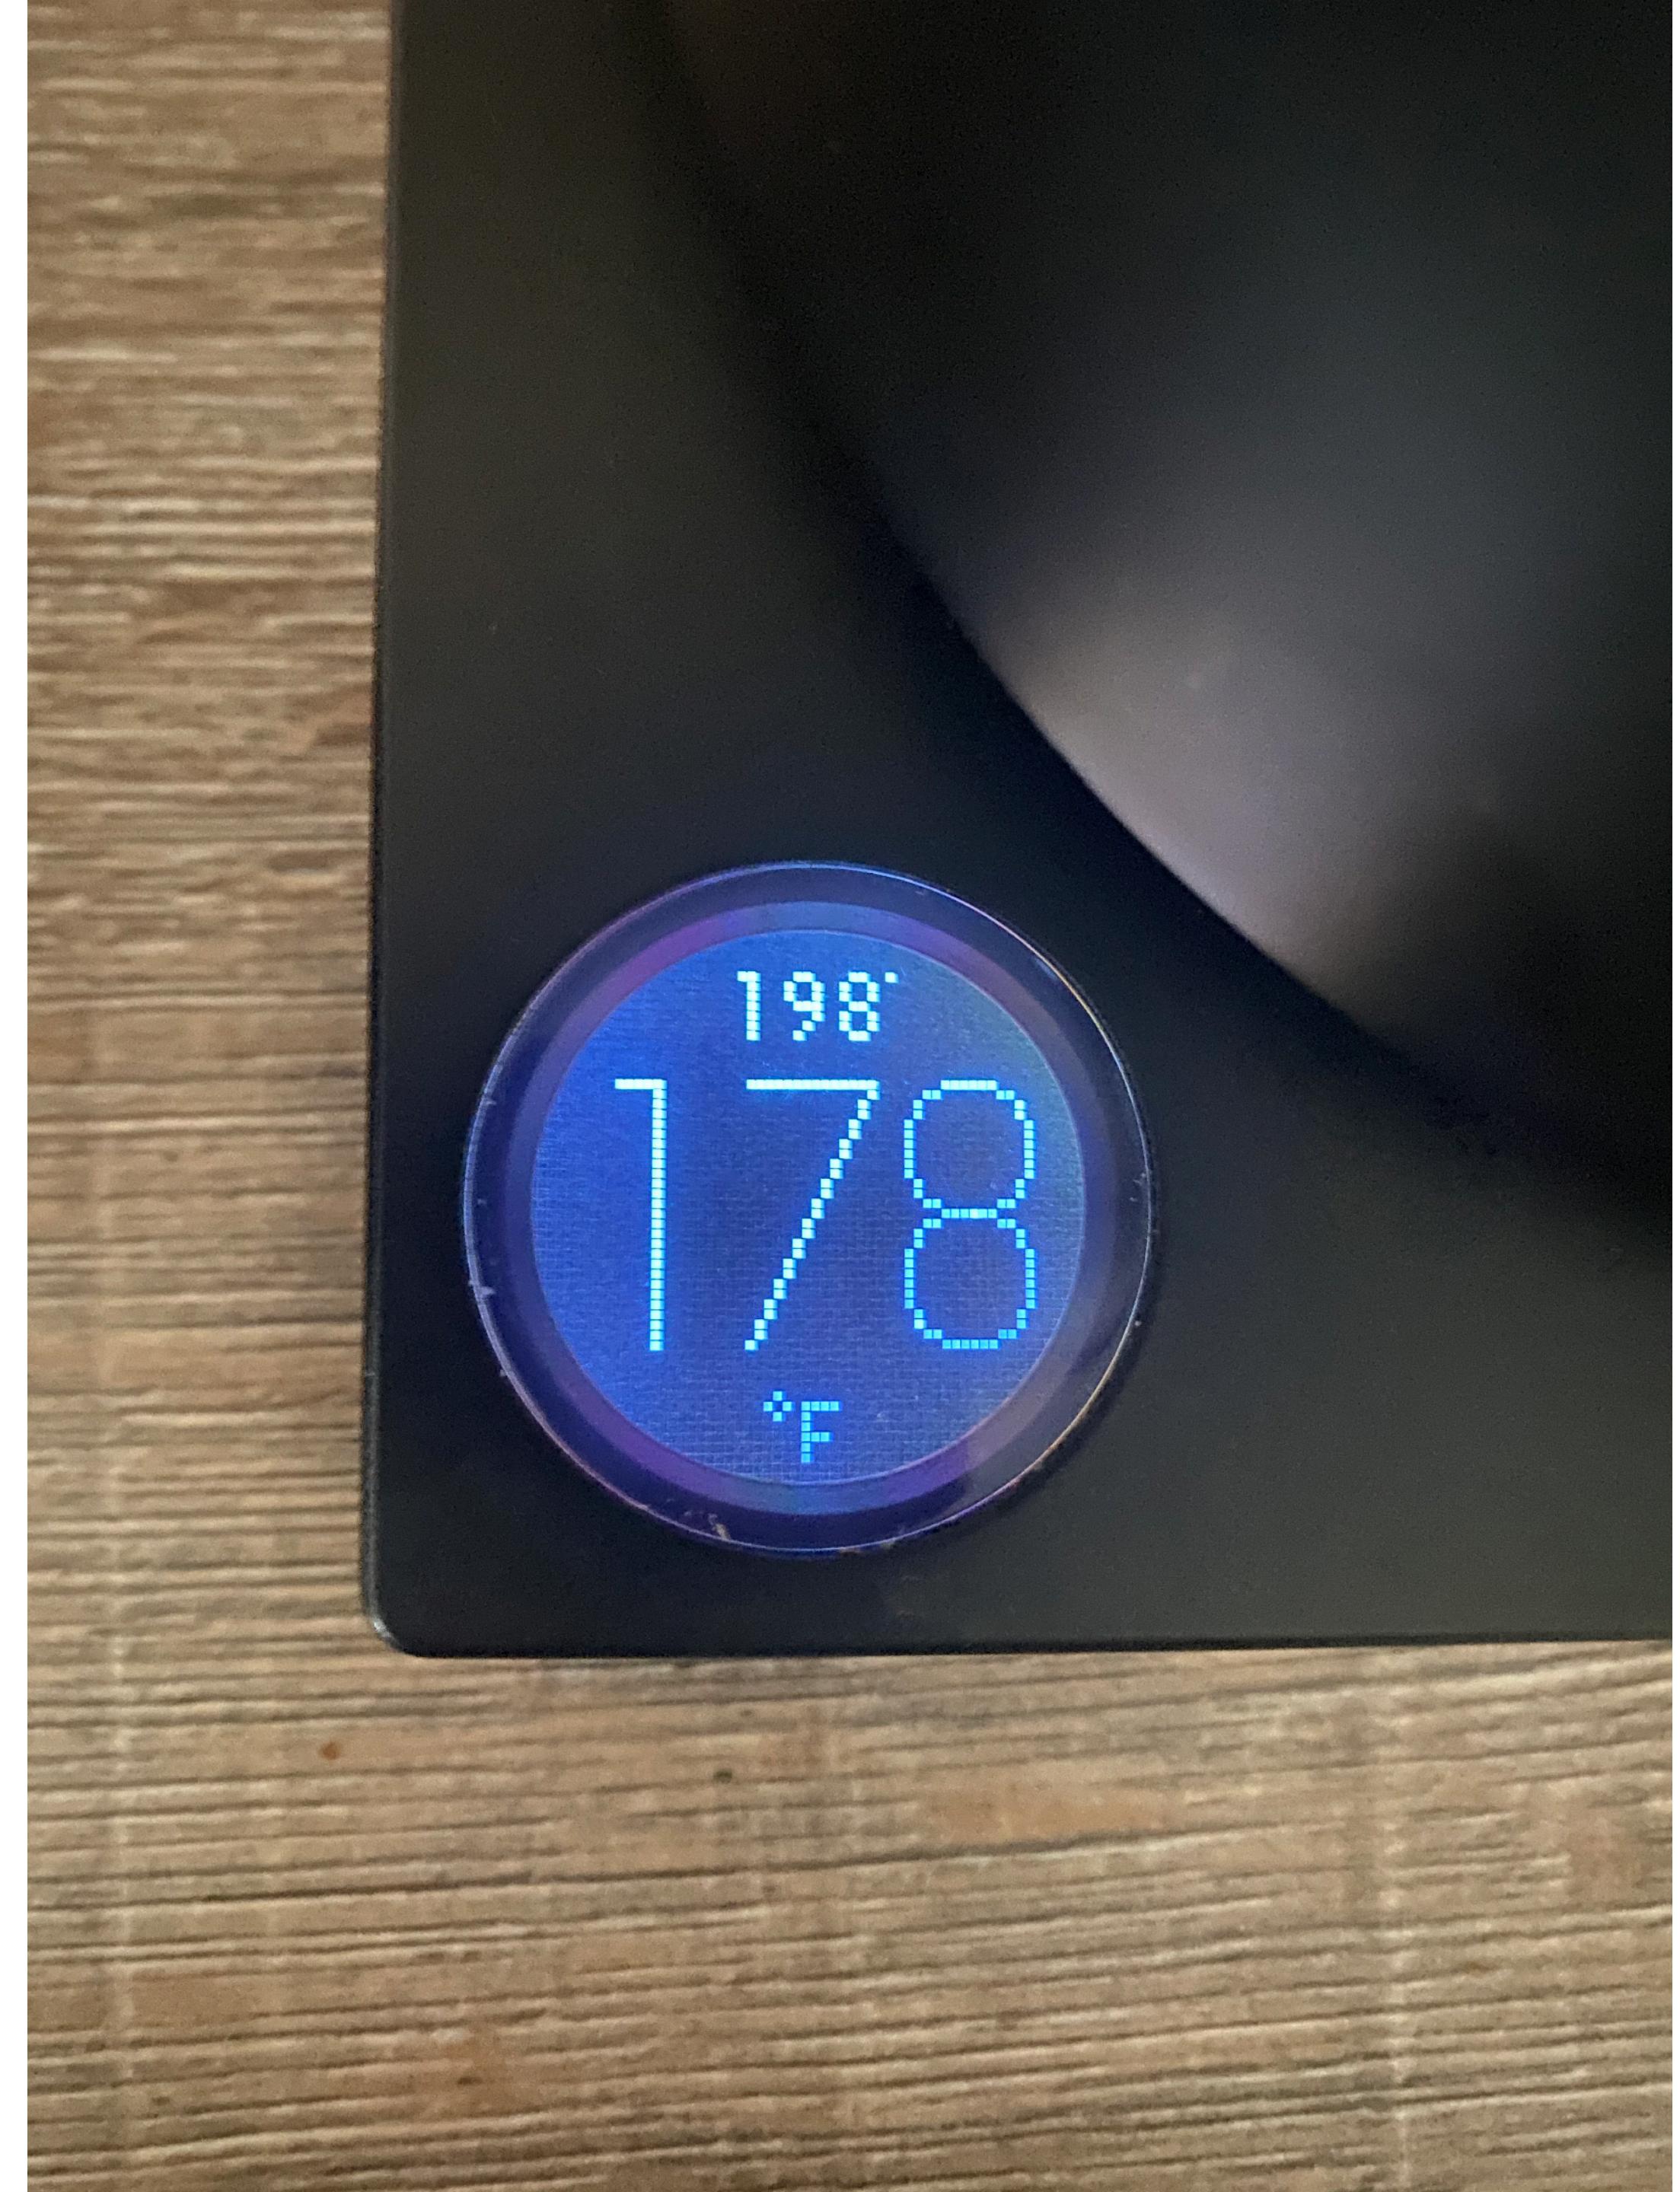
session.addInput(deviceInput)
let captureConnection = videoDataOutput.connection(with: .video)

let requestHandler = VNImageRequestHandler(cgImage: cgImage, options: [:])
let request = VNGenerateImageFeaturePrintRequest()
do {
    try requestHandler.perform([request])
    return request.results?.first as? VNFeaturePrintObservation
}
```

CODING 4: INTEGRATING VISION FRAMEWORK

VISION FRAMEWORK REQUESTS

- Follows the same GCD pattern from the legal app
- requires a **CGImage**
- Image types
 - **UIImage**: High Level Image for representation in UI (png/jpeg)
 - **CGImage**: Quartz CoreGraphic images, needs to be a bitmap
 - **CILImage**: Core Image which is basically data that can be transformed into an image



```
1 import UIKit
2 import Vision
3
4 let image = (UIImage(named: "image1.jpeg")?.cgImage)!

5
6 let imageRequestHandler = VNImageRequestHandler(cgImage: image, orientation: .right, options: [:])
7 let request = VNRrecognizeTextRequest { req, error in
8     guard let observations = req.results as? [VNRecognizedTextObservation] else {
9         return
10    }
11    let possibleResults: [Int] = observations.compactMap { observation in
12        var result: Int? = nil
13        for candidate in observation.topCandidates(15) {
14            guard let number = Int(candidate.string) else {
15                continue
16            }
17            result = number
18            break
19        }
20
21        return result
22    }
23
24    if possibleResults.count < 2 {
25        return
26    }
27    let current = possibleResults.min()!
28    let target = possibleResults.max()!

29
30    DispatchQueue.main.async {
31        print("Current temperature: \(current)")
32        print("Target temperature: \(target)")

33    }
34}
35 do {
36     try imageRequestHandler.perform([request])
37 } catch {
38     print("error occurred during processing")
39 }
40 }
```



```
Current temperature: 178
Target temperature: 198
```

WHERE TO GO FROM HERE

PART 4: DESIGN FOR PRIVACY

Approaches

What we didn't do in our Apps

Neural Engine

Importance of Privacy Now

Future-Proof Privacy

APPROACHES TO USING AI ON MOBILE DEVICES

- CLOUD APPROACH
- ON-DEVICE

WHAT WE DIDN'T DO IN OUR APPS

- We didn't store the images in our cloud
- We didn't process the images in our via REST service for our AI in the cloud
- Pros
 - we don't pay for storage for the images
 - we don't have costs for the API calls and CPU/GPU resources
 - we don't have service outages or maintenance costs
- Cons
 - We cannot improve our AIs based on the images from users
 - It's harder to argue for subscription costs based on data storage & api calls
 - Some AI tasks may be too complex for the phone's neural engine to handle

NEURAL ENGINE

- Using the AIs on device is only possible with a specialized hardware that can run AIs efficiently
- The Power of The Neural Engine
- Ensures Privacy and Reduces our deployment costs
- However our AIs, need to be adjusted so that they are best compatible with the Neural Engine
- We need to make sure that we don't drain the battery and still have a responsive app

THE IMPORTANCE OF PRIVACY - NOW

- **Self-fulfilling prophecy: Some people actually care about privacy (even in the 🇺🇸)**
 - These users might even be skewed towards Apple's ecosystem - target them!
- **Designing for privacy means NOT HAVING TO TAKE responsibility for**
 - Data Leaks / Hacker attacks
 - Data Management / Big Data (do you want to end up with Hadoop?)
 - Sending stuff back and forth
 - Your users' privacy!
- **By valuing privacy now, you build trust for your brand later!**

THE IMPORTANCE OF PRIVACY IN THE FUTURE

- Let's imagine Apple comes up with something you wear on / in your eyes
 - It is permanently recording - everything
 - Eventually, nobody would want to have everything they see shared with profit seeking companies
 - Apple might not allow it, laws might prohibit it!
- There might be a future with body parts that constantly process data which you most definitely don't want to share
 - In-Body Device DNA/RNA Testing
- The future does not have to be a dystopia
- Privacy preserving intelligence might become mandatory.